

Is my advertising working?

Four methods for measuring advertising response

Elea McDonnell Feit

10 May 2017

updated 1 September 2017

Topics

Preliminaries

Attribution rules

Holdout testing

Marketing mix modeling

Model-based attribution

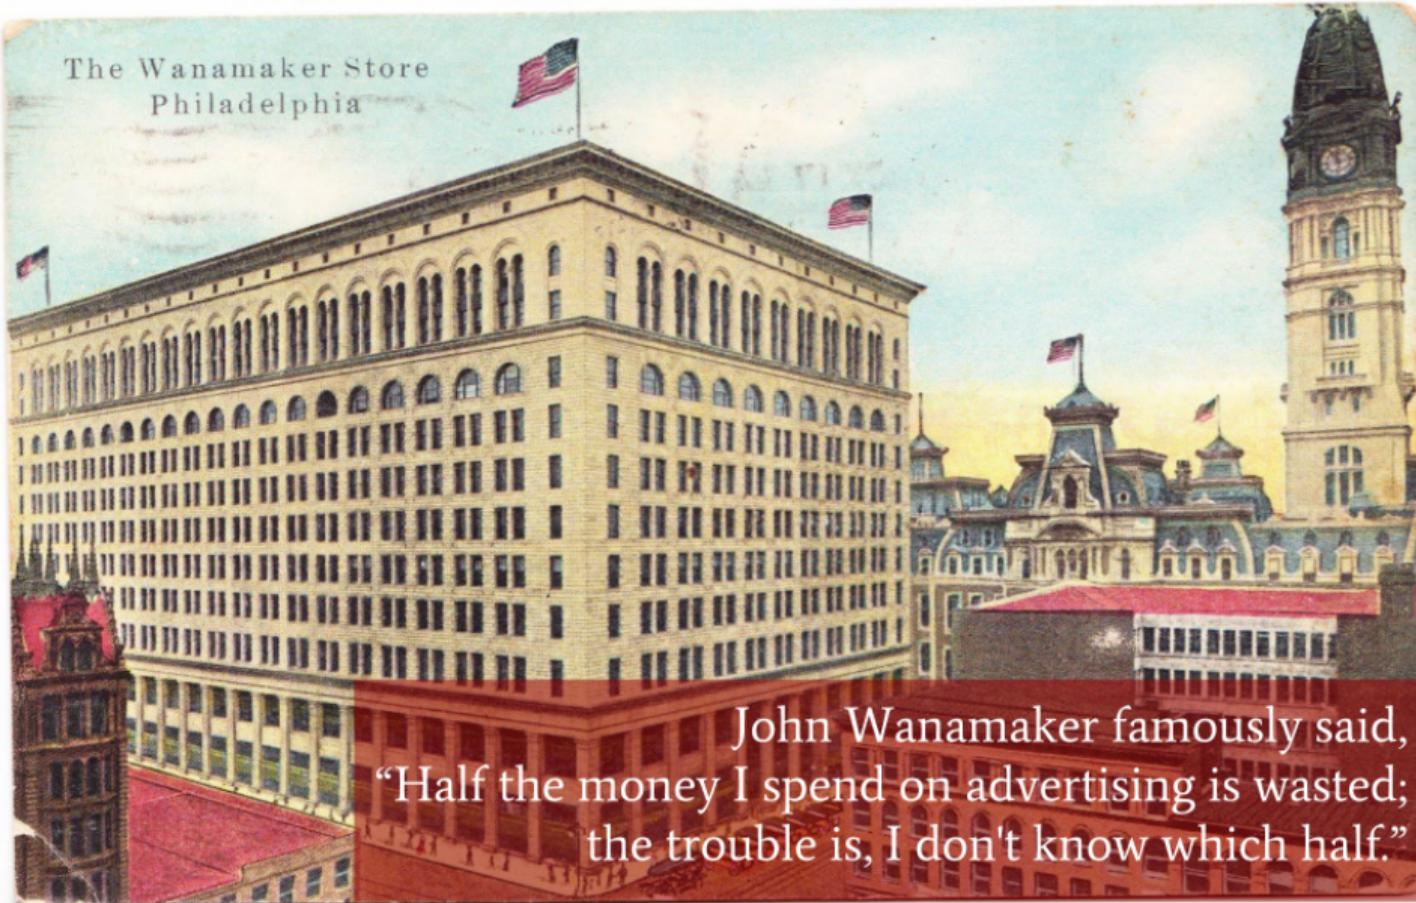
Review

Preliminaries

Preliminaries

Why advertising response?

The Wanamaker Store
Philadelphia



John Wanamaker famously said,
“Half the money I spend on advertising is wasted;
the trouble is, I don't know which half.”

Measuring advertising response

The goal of any marketing campaign is to increase sales (short-term or long-term).

Measuring advertising response

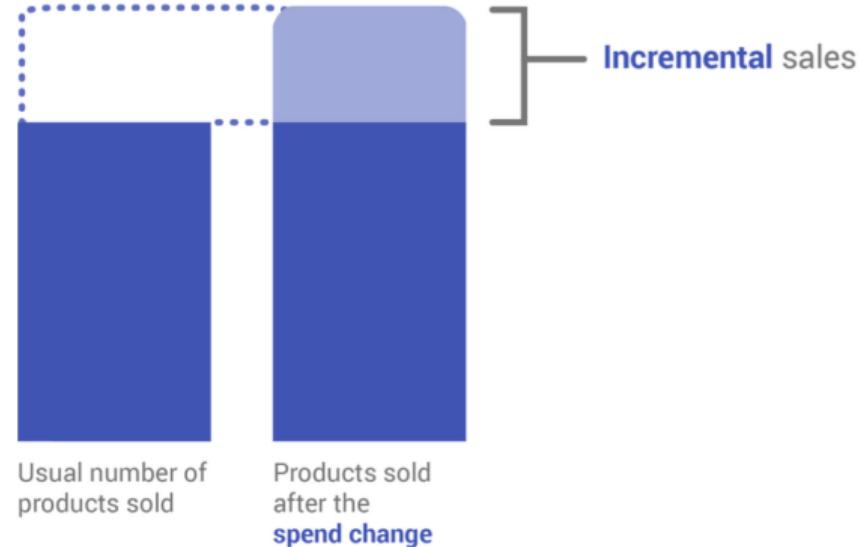
The goal of any marketing campaign is to increase sales (short-term or long-term).

In theory, it should be easy to evaluate the performance of marketing. Each campaign or marketing channel should be evaluated based on the **incremental profit** that it produces relative to its **cost**.

$$ROI = \frac{\text{incremental profit due to advertising} - \text{cost of advertising}}{\text{cost of advertising}}$$

Incremental sales

Incremental profit depends on **incremental sales** which are the additional sales we make with advertising over and above what we would have sold without advertising. Incremental profit is typically a function of the incremental sales.

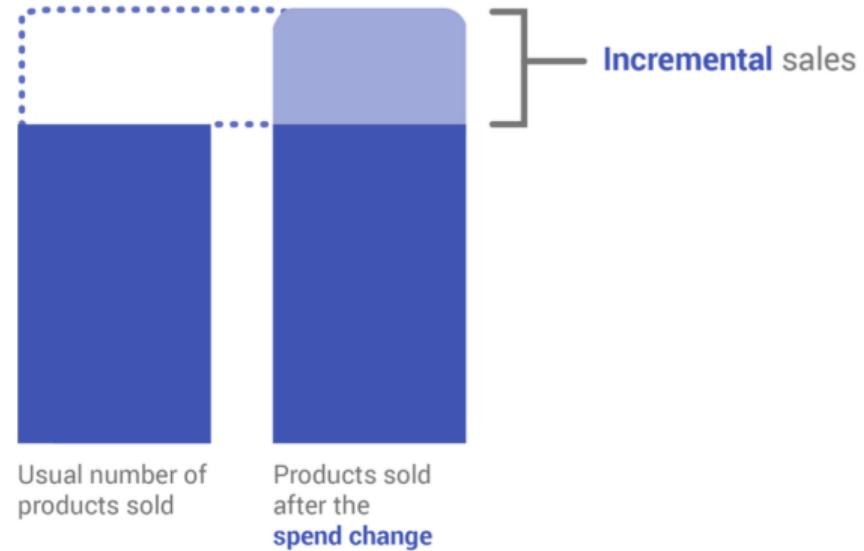


Source: [Think with Google](#)

Incremental sales

Incremental profit depends on **incremental sales** which are the additional sales we make with advertising over and above what we would have sold without advertising. Incremental profit is typically a function of the incremental sales.

With today's digital media – where we can track individual users – it is now possible to estimate incremental sales in many situations.



Source: [Think with Google](#)

Workshop goals

Understand four popular techniques for estimating the incremental sales due to advertising.

- Attribution rules
- Holdout testing
- Marketing mix modeling (with aggregate spending and sales data)
- Attribution modeling (with user-level data)

Workshop goals

Understand four popular techniques for estimating the incremental sales due to advertising.

- Attribution rules
- Holdout testing
- Marketing mix modeling (with aggregate spending and sales data)
- Attribution modeling (with user-level data)

Learn how each technique works using the R statistical language.

Workshop goals

Understand four popular techniques for estimating the incremental sales due to advertising.

- Attribution rules
- Holdout testing
- Marketing mix modeling (with aggregate spending and sales data)
- Attribution modeling (with user-level data)

Learn how each technique works using the R statistical language.

Understand which techniques produce better estimates of incremental profits and why.

Workshop goals

Understand four popular techniques for estimating the incremental sales due to advertising.

- Attribution rules
- Holdout testing
- Marketing mix modeling (with aggregate spending and sales data)
- Attribution modeling (with user-level data)

Learn how each technique works using the R statistical language.

Understand which techniques produce better estimates of incremental profits and why.

After this workshop, you won't be an expert, but you will have a much better idea of how these techniques work and you will be a much better consumer of attribution analysis that is done for you.

What if I don't know R?

I will assume that you have some familiarity with regression, but I don't expect everyone here knows R.

What if I don't know R?

I will assume that you have some familiarity with regression, but I don't expect everyone here knows R.

There are two ways to follow this tutorial:

What if I don't know R?

I will assume that you have some familiarity with regression, but I don't expect everyone here knows R.

There are two ways to follow this tutorial:

What if I don't know R?

I will assume that you have some familiarity with regression, but I don't expect everyone here knows R.

There are two ways to follow this tutorial:

1. "Go for a ride"

- Focus on the data that goes into the analysis and the output of the analysis, and how we interpret it.
- Don't worry too much about the R syntax. Let me "drive" that.
- Review the R code later to learn the syntax (if you want).

What if I don't know R?

I will assume that you have some familiarity with regression, but I don't expect everyone here knows R.

There are two ways to follow this tutorial:

1. "Go for a ride"

- Focus on the data that goes into the analysis and the output of the analysis, and how we interpret it.
- Don't worry too much about the R syntax. Let me "drive" that.
- Review the R code later to learn the syntax (if you want).

2. "Practice driving"

- For those comfortable using R and to do data manipulation and regression (all base R, no tidyverse), download the code file and follow along.

What if I don't know R?

I will assume that you have some familiarity with regression, but I don't expect everyone here knows R.

There are two ways to follow this tutorial:

1. "Go for a ride"

- Focus on the data that goes into the analysis and the output of the analysis, and how we interpret it.
- Don't worry too much about the R syntax. Let me "drive" that.
- Review the R code later to learn the syntax (if you want).

2. "Practice driving"

- For those comfortable using R and to do data manipulation and regression (all base R, no tidyverse), download the code file and follow along.

I don't intend to teach you R. I'm just using because it is my favorite platform for data manipulation, computing summary statistics, graphing and regression modeling. This analysis could be done using Python/Pandas, SAS, Matlab, etc.

Preliminaries

Data

Data: retail impressions and conversions

Throughout this workshop, we will be analyzing a data set describing 10,000 customers and potential customers of a retailer. The retailer uses four different advertising channels.

Display

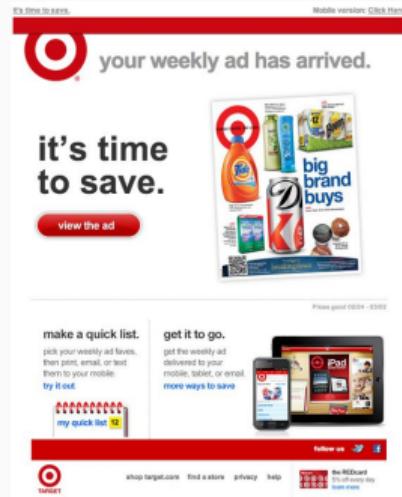


Social

Sponsored

A social media post from 'Target Style target.com'. It shows a woman wearing a floral skirt and a tan wedge sandal. The post includes the text 'Target Style target.com' and 'Love these? Make 2017 your most stylish year yet.' There's also a 'Create Ad' button.

Email



Direct



Customer tracking begins when the customer is exposed to a display or social ad, visits the retailer website or makes a purchase.

Data structure

The data is organized in three files:

- `customer.csv`: each row is a customer, 10,000 rows
- `impressions.csv`: each row is an exposure of a marketing communication to a specific customer, 501,336 rows
- `transactions.csv`: each row is a transaction made by a customer

Data structure

The data is organized in three files:

- `customer.csv`: each row is a customer, 10,000 rows
- `impressions.csv`: each row is an exposure of a marketing communication to a specific customer, 501,336 rows
- `transactions.csv`: each row is a transaction made by a customer

This basic structure is typical of raw digital advertising data. The data is hypothetical – I made it in R – and has been designed to illustrate important points in data analysis.

Some R preliminaries

We start our R session with a bit of tidying up.

```
> rm(list=ls()) # clears the workspace
> set.seed("20030603") # ensures repeatable results for attribution rules
> options(scipen=999) # suppress scientific notation
```

Customer file

The customer file can be downloaded from goo.gl/fyxRE5. Each row describes a customer.

Columns:

id: an id number for the customer

past purchase: whether the customer has made a purchase prior to the observation period.

email: indicates whether the customer is eligible to receive emails, i.e. we have an email address and permission to mail

direct: indicates whether the customer is eligible to receive direct mail, i.e. we have an address and permission to mail

Read and inspect customer file

We can read the file directly into R:

```
> cust <- read.csv("https://goo.gl/mqy8NR")
> nrow(cust)
[1] 10000
> summary(cust)
      id      past.purchase      email      direct
Min.   : 1   Min.   :0.0000   Min.   :0.0000   Min.   :0.0000
1st Qu.: 2501 1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.:0.0000
Median : 5000 Median :1.0000   Median :1.0000   Median :0.0000
Mean   : 5000 Mean   :0.5022   Mean   :0.6001   Mean   :0.4974
3rd Qu.: 7500 3rd Qu.:1.0000   3rd Qu.:1.0000   3rd Qu.:1.0000
Max.   :10000 Max.   :1.0000   Max.   :1.0000   Max.   :1.0000
```

We have 10,000 customers and about half of those have made a purchase. About 60% are eligible for email and 50% are eligible for direct mail.

More inspection of customer file

```
> xtabs(~email+direct+past.purchase , data=cust)
, , past.purchase = 0
  direct
email      0      1
  0 2775  707
  1 1205  291
, , past.purchase = 1
  direct
email      0      1
  0    95   422
  1  951  3554
```

Most customers who have not made a purchase are not eligible for email or direct mail.
Customers who have made a purchase are more likely to be eligible.

Impressions file

The impressions file can be downloaded from <https://goo.gl/LT93AA>. Each row in the file represents an exposure of one customer to an ad, i.e. an impression.

Columns:

id: id number for the customer

date: date of impression. Most files would have a date-time, but we have simplified for the workshop.

channel: channel of the ad exposure

click: indicates whether the customer clicked on the ad

Read and inspect impressions file

```
> impress <- read.csv("https://goo.gl/74qIxY")
> impress$date <- as.Date(impress$date) # change type
> nrow(impress)
[1] 501336
> summary(impress)
      id           date          channel        click
Min.   : 1   Min.   :2016-12-31   direct     : 9948   Min.   :0.00000
1st Qu.: 2467 1st Qu.:2017-01-10  display    :216371  1st Qu.:0.00000
Median : 4940 Median  :2017-01-20  email     : 38426  Median :0.00000
Mean   : 4960 Mean   :2017-01-22  email.holdout: 9582   Mean   :0.01854
3rd Qu.: 7454 3rd Qu.:2017-01-31  social    :227009  3rd Qu.:0.00000
Max.   :10000  Max.   :2017-02-27
```

We have 501,336 impressions between 2016-12-31 to 2017-02-27. The majority of these impressions are social or display. Overall click rate is 1.85%, which is reasonable.

Summarize the cadence of the impressions

Cadence is the timing of advertising impressions. It is useful to understand the cadence of your ads before you do any other analysis. Email and direct are often sent out to users on specific dates. Display and social ads can be steady or can be turned on and off at specific times. We can summarize the cadence with the crosstab function in R.

```
> (cadence <- xtabs(~date+channel, data=impress))
```

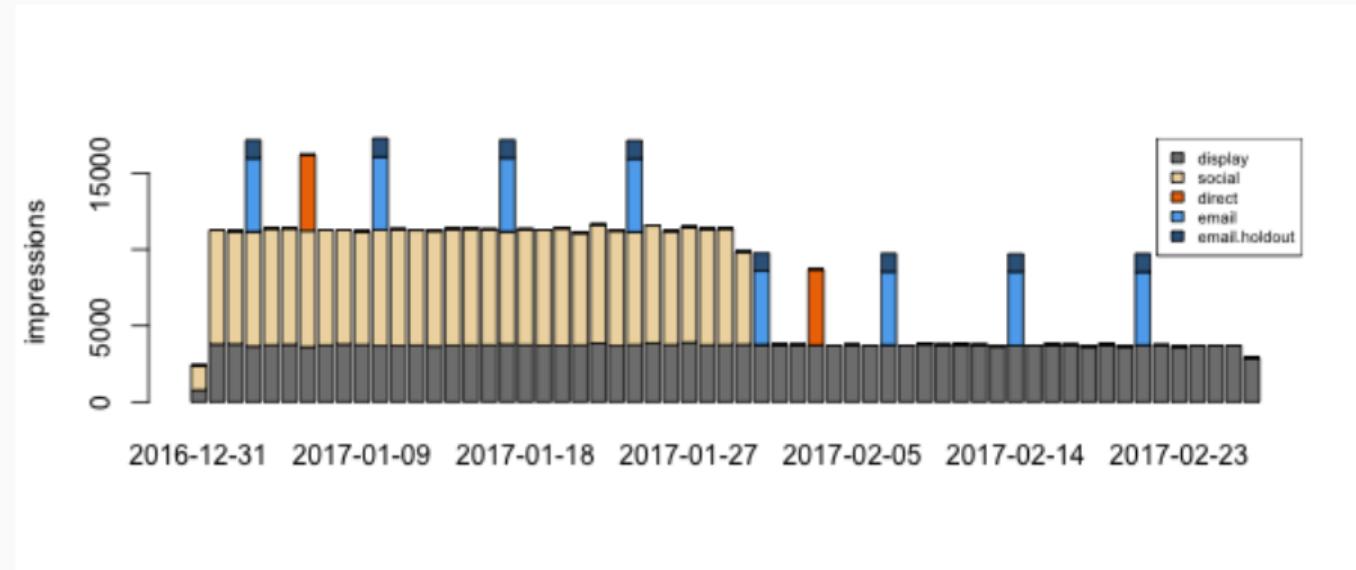
date	channel				
	direct	display	email	email.holdout	social
2016-12-31	0	788	0	0	1610
2017-01-01	0	3786	0	0	7481
2017-01-02	0	3792	0	0	7416
2017-01-03	0	3656	4798	1203	7505
2017-01-04	0	3731	0	0	7648
...					

Visualizing the cadence

It is useful to create plot of the cadence data.

```
> cadence <- cadence[,c(2,5,1,3,4)] # reorder the columns  
> mycols <- c("gray50", "wheat2", "darkorange2", "steelblue2", "steelblue4")  
> barplot(t(cadence), col=mycols, ylab="impressions")  
> legend("topright", legend=colnames(cadence), fill=mycols, cex=0.6)
```

Visualizing the cadence



Display impressions per day are steady across the observation window.

Social is steady in the first month and then stops.

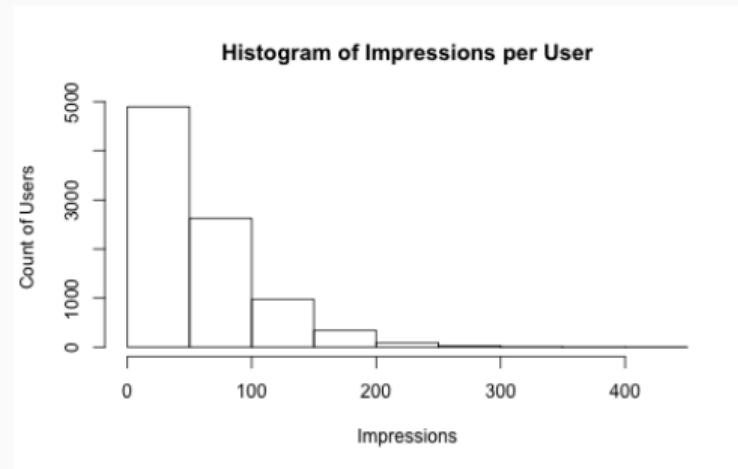
Emails are sent once per week and each campaign seems to have a holdout.

Direct mail is sent once per month.

Histogram of impressions

It is also useful to understand how many impressions each customer gets.

```
> hist(xtabs(~id, data=impress), xlab="Impressions", ylab="Count of Users",  
+       main="Histogram of Impressions per User")
```



Some customers receive as many as 400 impressions in two months, but most get less than 100 impressions.

Click through rates by channel

```
> xtabs(click~channel, data=impress)/xtabs(~channel, data=impress)
channel
    direct      display      email  email.holdout      social
0.0000000000  0.004783451  0.099672097  0.000000000  0.019519050
```

The click through rates are highest for email at 10.0% and lowest for display at 0.5%. As expected, there are no clicks for direct or email holdouts.

Transactions file

The transactions file can be downloaded from <https://goo.gl/ZErZ6g>. Each row in the file represents a purchase made by a customer.

Columns:

`id`: id number for the customer

`date`: date of transactions. Most files would have a date-time, but we have simplified for the workshop.

`last.touch`: channel of the last ad impression the customer saw before the transaction

`last.click`: channel of the last ad the customer clicked before the transaction

Purchase amounts can be used in holdout testing, attribution models or marketing mix models, but we will focus on transaction counts for this workshop.

Read and inspect transactions file

```
> summary(trans)
```

	X	id	date	last.touch	last.click
Min.	: 1	Min. : 2	Min. :2017-01-01	direct :2594	display:
891					
1st Qu.	: 5609	1st Qu.: 2472	1st Qu.:2017-01-12	display :5252	email : 4151
Median	:11217	Median : 5005	Median :2017-01-25	email : 7145	none :12998
Mean	:11217	Mean : 4985	Mean :2017-01-26	none : 846	social : 4393
3rd Qu.	:16825	3rd Qu.: 7483	3rd Qu.:2017-02-08	social :6596	
Max.	:22433	Max. :10000	Max. :2017-02-28		

The transactions span the same dates and customer ids as the impressions. It looks like some customers such as id 1 do not have transactions.

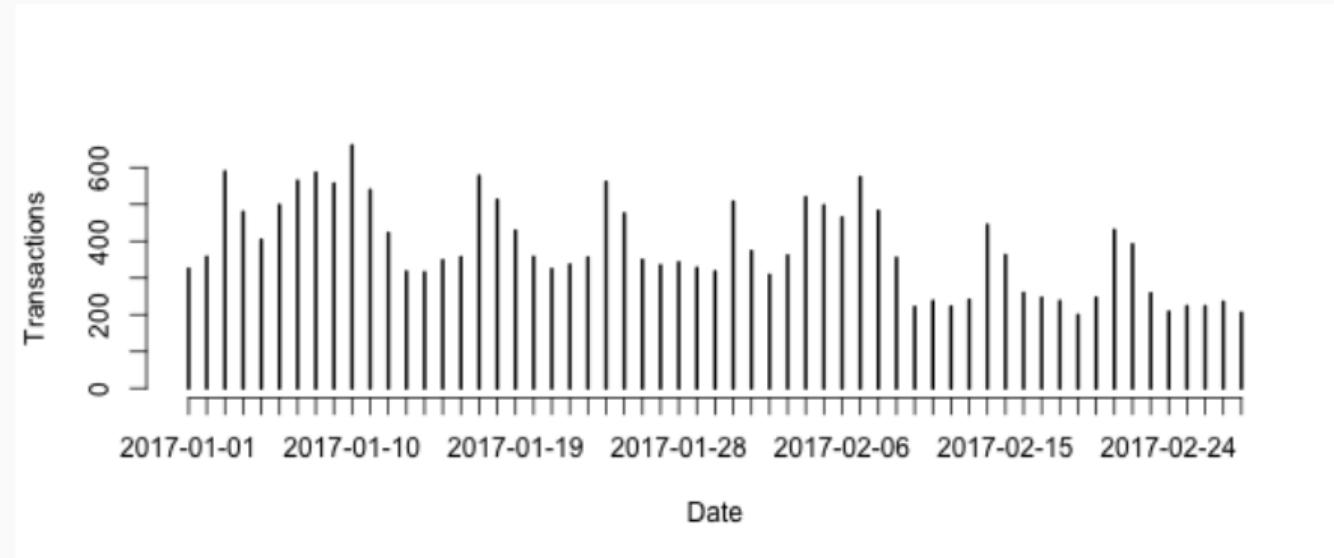
Transactions over time

```
> (transbyday <- xtabs(~date, data=trans))
date
2017-01-01 2017-01-02 2017-01-03 2017-01-04 2017-01-05 2017-01-06 2017-01-07 2017-01-08
      325        357        589        479        403        498        564        586
2017-01-09 2017-01-10 2017-01-11 2017-01-12 2017-01-13 2017-01-14 2017-01-15 2017-01-16
      556        660        539        422        317        315        348        356
2017-01-17 2017-01-18 2017-01-19 2017-01-20 2017-01-21 2017-01-22 2017-01-23 2017-01-24
      577        512        428        357        324        336        355        561
2017-01-25 2017-01-26 2017-01-27 2017-01-28 2017-01-29 2017-01-30 2017-01-31 2017-02-01
      475        349        334        342        327        317        507        373
2017-02-02 2017-02-03 2017-02-04 2017-02-05 2017-02-06 2017-02-07 2017-02-08 2017-02-09
      ...
...
```

This is easier to understand if we plot it.

Plot of transactions over time

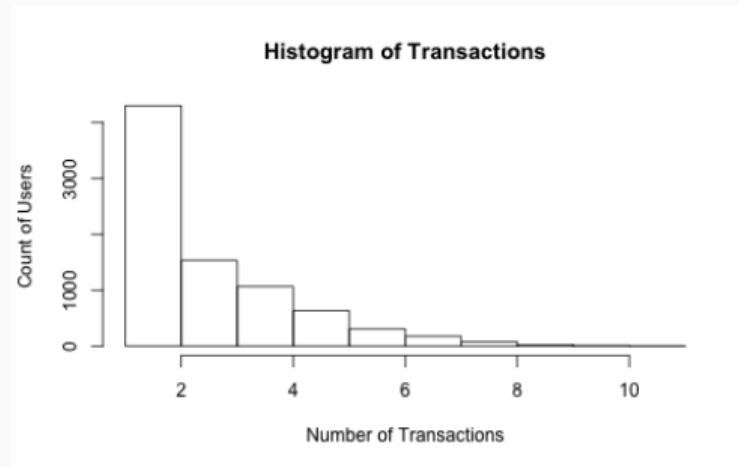
```
> plot(transbyday, ylab="Transactions", xlab="Date")
```



Transactions appear to be a bit higher in the first month and there are clear spikes around the time of emails and direct mail.

Histogram of transactions

```
> hist(xtabs(~id, data=trans), xlab="Number of Transactions", ylab="Count of Users",
+       main="Histogram of Transactions")
```



Most customers have less than 2 transactions in the two month observation period.

Customer 100

To understand how the files fit together, it is helpful to look at the impressions and transactions for one customer.

```
> cust[cust$id==100,]
  id past.purchase email direct
100 100            0    0    0
> impress[impress$id==100,]
[1] id      date    channel click
<0 rows> (or 0-length row.names)
> trans[trans$id==100,]
  X   id      date last.touch last.click
206 206 100 2017-01-18       none       none
207 207 100 2017-01-26       none       none
```

Customer 100 has no impressions and made two transactions.

Customer 300

```
> cust[cust$id==300,]
      id past.purchase email direct
300 300           1     1     1
> summary(impress[impress$id==300,])
      id          date            channel      click
Min. :300  Min.   :2016-12-31  direct       : 2  Min.   :0
1st Qu.:300  1st Qu.:2017-01-10  display      :81  1st Qu.:0
Median :300  Median  :2017-01-29  email        : 8  Median  :0
Mean   :300  Mean    :2017-01-28  email.holdout: 0  Mean   :0
3rd Qu.:300  3rd Qu.:2017-02-14  social       : 0  3rd Qu.:0
Max.   :300  Max.    :2017-02-26                    Max.   :0
> trans[trans$id==300,]
      X id          date last.touch last.click
647 647 300 2017-01-03   display      none
648 648 300 2017-01-10   email        none
649 649 300 2017-02-04   direct       none
650 650 300 2017-02-19   display      none
```

Customer 300 has 8 email, 2 direct mail and 81 display impressions and made 4 transactions.

Steps in data inspection

1. For each file,
 - 1.1 Make sure you have the entire file read in by checking number of rows.
 - 1.2 Summarize individual variables in file mean, min, max, etc. Do they make sense?
 - 1.3 Summarize relationships between variables in file using crosstabs or scatterplots. Do they make sense?
2. Check that joins between files are as expected (skipped).
3. Summarize relationships between files. Do they make sense?

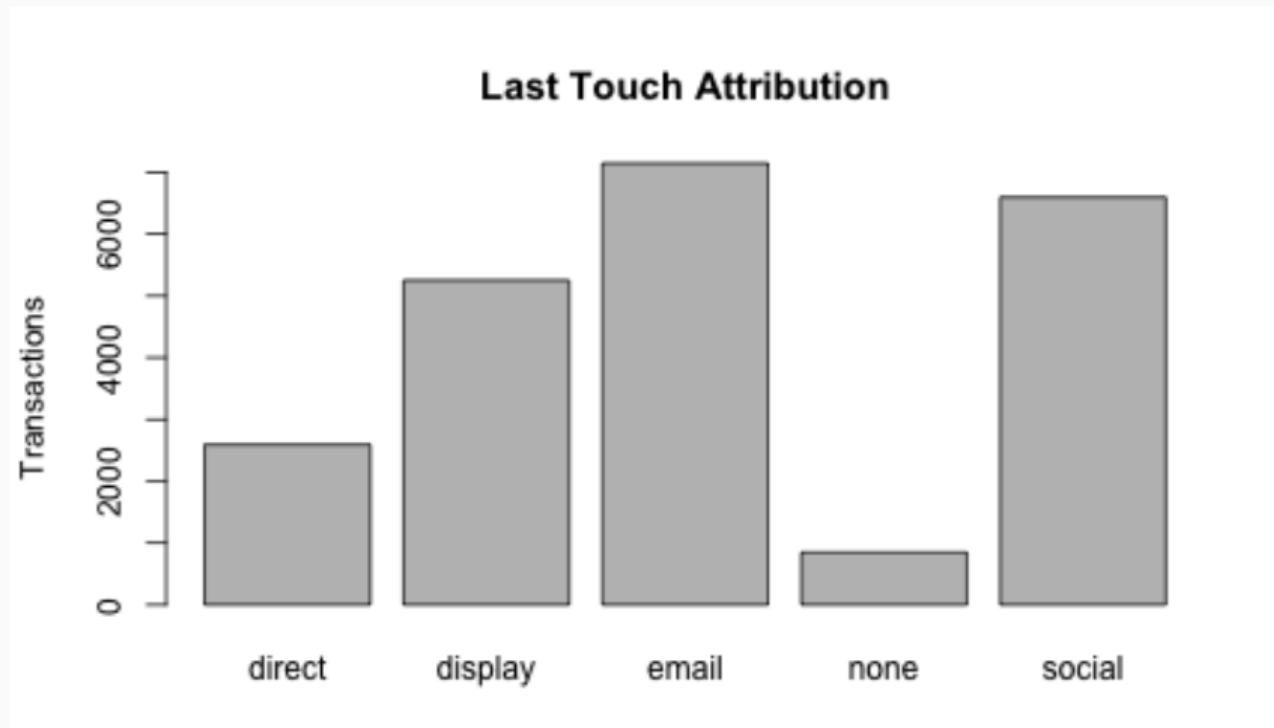
Attribution rules

Attribution rules

What is last click?

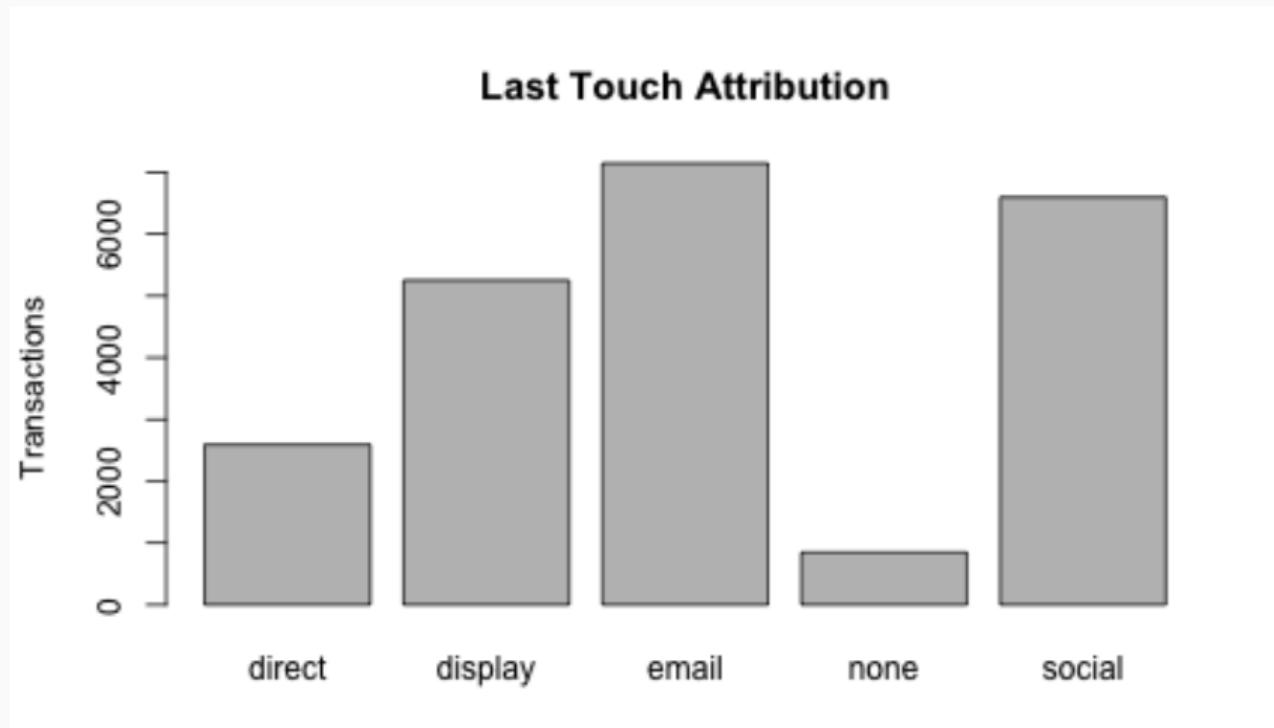
Last-touch attribution

Last-click (or last-touch) attribution looks backward from each conversion to find the last ad the user clicked on prior to the conversion.



Last-touch attribution

Last-click (or last-touch) attribution looks backward from each conversion to find the last ad the user clicked on prior to the conversion.



Attribution rules

Last touch analysis in R

Last touch analysis

The last click and the last touch are stored in the transaction file.

```
> head(trans)
  id      date last.touch last.click
1  2 2017-01-04      email      none
2  2 2017-02-12      email      none
3  3 2017-02-02      email      none
4  3 2017-02-14      email      none
5  5 2017-01-04    display      email
6  5 2017-01-13    display      email
```

One advantage of last-click is that it is easy to compute once this pre-processing is done.

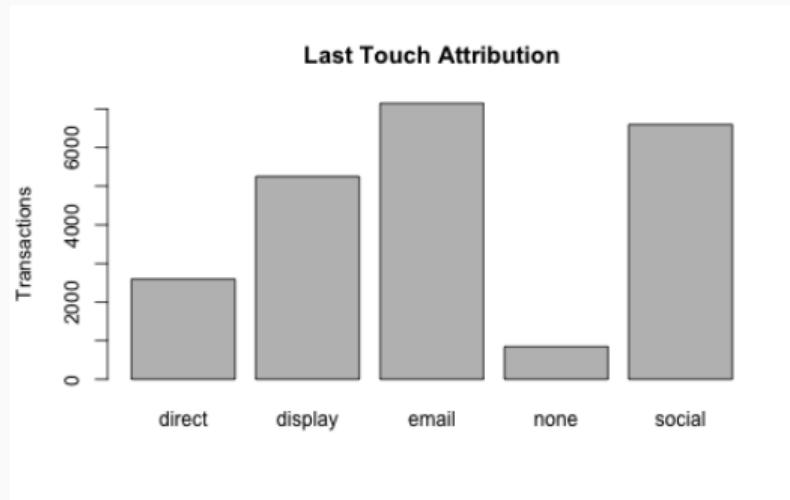
Last touch analysis

For example, if we want to compute the number of transactions that are attributed to each channel by last touch, we just do a quick crosstab on the transaction table.

```
> (last.touch.att <- xtabs(~last.touch, data=trans))
last.touch
direct display   email     none   social
    2594      5252     7145      846     6596
> barplot(last.touch.att, ylab="Transactions",
+           main="Last Touch Attribution")
```

When we do this, we are ignoring all the customers who didn't transact.

Last touch analysis



Many people interpret this as meaning that the incremental sales for social are 6596.

Last touch for subgroups of transactions

It is also easy to compute the last touch for a subgroup of the transactions. Simply subset out the transactions of interest and then crosstab.

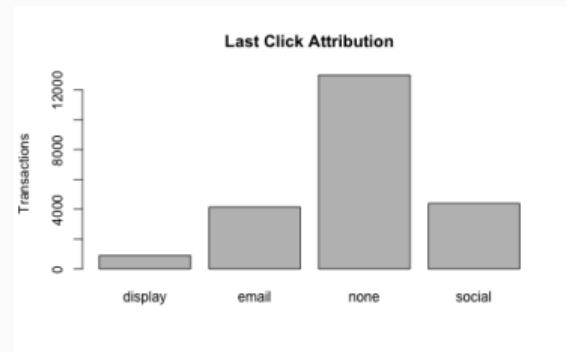
```
> xtabs(~last.touch, data=trans[trans$date>as.Date("2017-01-31"),])  
last.touch  
direct display email none social  
1669    2682    4081    258    328
```

In February, there were far fewer sales attributed to social, most likely because we ended the social ads at the end of January.

Last click analysis

You can do the same analysis based on the last ad a customer clicked, rather than the last touch/impression.

```
> (last.click.att <- xtabs(~last.click, data=trans))
last.click
display   email     none   social
      891     4151    12998    4393
> barplot(last.click.att, ylab="Transactions",
+           main="Last Click Attribution")
```



Since people don't click very much, the attributed sales are much lower. Measuring ad performance based on clicks doesn't make much sense for the advertiser.

Appending last click data [bonus R material]

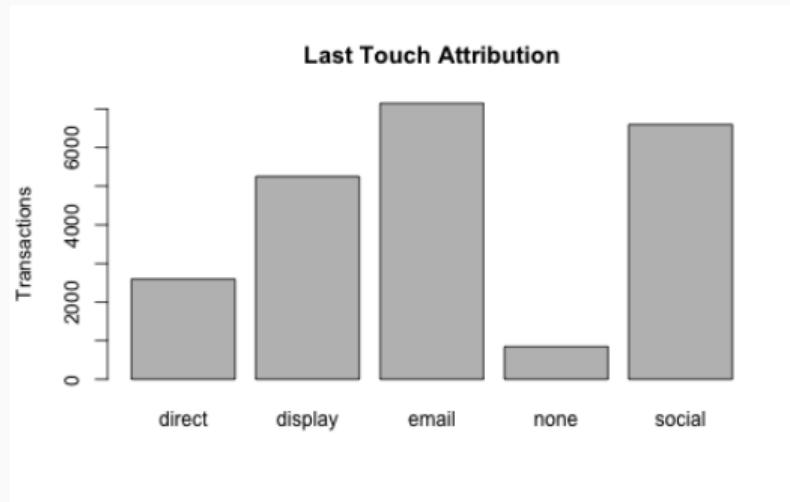
This function takes a transaction table that does not include the last touch and appends it based on the impression table.

```
append.last <- function(trans, impress) {  
  out <- data.frame(trans, last.touch=NA, last.click=NA)  
  for (t in 1:nrow(out)) {  
    impr <- impress[impress$id==out$id[t] & impress$date<=out$date[t] & impress$channel!="email.holdout", ]  
    if (nrow(impr)>0) {  
      out$last.touch[t] <- as.character(sample(impr$channel[impr$date==max(impr$date)], 1)) # choose randomly  
    }  
    impr <- impr[impr$click==1,]  
    if (nrow(impr)>0){  
      out$last.click[t] <- as.character(sample(impr$channel[impr$date==max(impr$date)], 1))  
    }  
  }  
  out[is.na(out)] <- "none"  
  out  
}  
trans <- append.last(trans, impress)
```

Attribution rules

Limitations of last-click

Last touch analysis



Many people interpret this as meaning that the incremental sales for social are 6596.

Incorrect assumptions behind last touch

When we use last touch to estimate the incremental sales for an ad, we are making two mistakes:

1. Other ads may have influenced the customer and contributed to the sale.

Incorrect assumptions behind last touch

When we use last touch to estimate the incremental sales for an ad, we are making two mistakes:

1. Other ads may have influenced the customer and contributed to the sale.
 - This is the problem everyone seems to recognize.

Incorrect assumptions behind last touch

When we use last touch to estimate the incremental sales for an ad, we are making two mistakes:

1. Other ads may have influenced the customer and contributed to the sale.
 - This is the problem everyone seems to recognize.
2. It counts all the sales as incremental, i.e. it assumes that customers who saw ads would not have bought if they hadn't seen the ads.

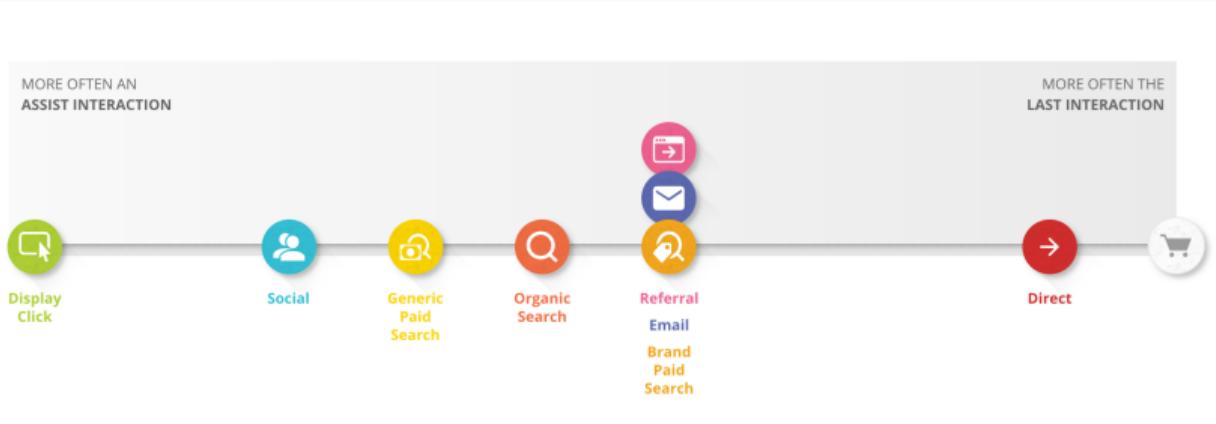
Incorrect assumptions behind last touch

When we use last touch to estimate the incremental sales for an ad, we are making two mistakes:

1. Other ads may have influenced the customer and contributed to the sale.
 - This is the problem everyone seems to recognize.
2. It counts all the sales as incremental, i.e. it assumes that customers who saw ads would not have bought if they hadn't seen the ads.
 - This issue is less well recognized.

Attribution as a payment mechanism

Because last-touch is so simple to compute, it has been used as a rule for allocating payments between channels on pay-for-conversion ads.

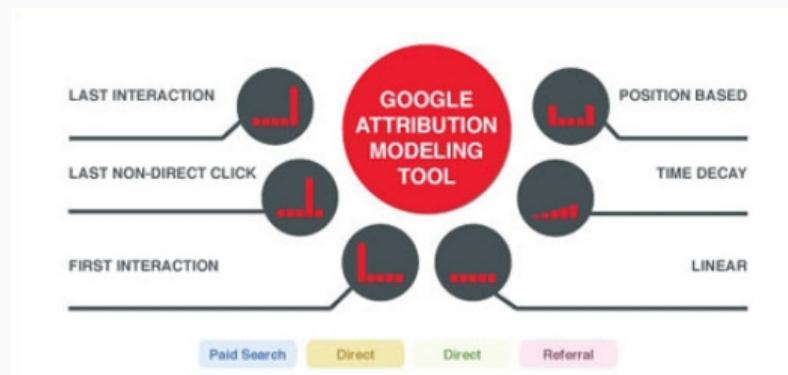


Source: [Think with Google](#)

Last touch unfairly favors channels that tend to show ads towards the end of the path to purchase such as search and retargeting.

Other attribution rules (“models”)

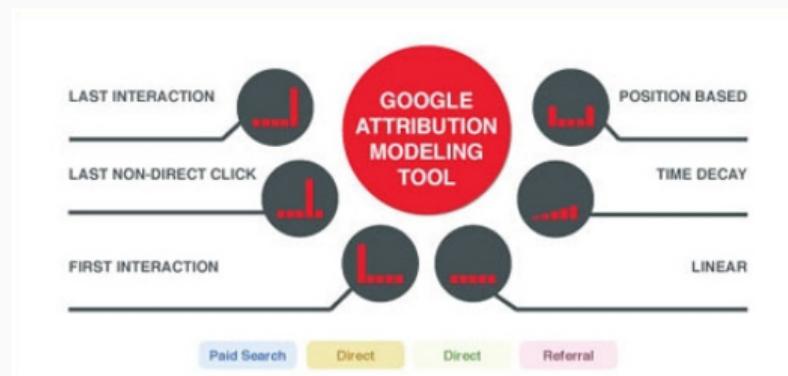
Perhaps because everyone intuitively understood that last touch didn't account for “assists” and was unfair to channels that tend to occur early in the path to purchase they came up with other arbitrary rules.



Source: [Across Health](#)

Other attribution rules (“models”)

Perhaps because everyone intuitively understood that last touch didn't account for “assists” and was unfair to channels that tend to occur early in the path to purchase they came up with other arbitrary rules.



Source: [Across Health](#)

These improved rules may help problem 1, but they don't help with problem 2.

When does last-touch “work”

Last click provides an accurate measure of ad response when

- All sales from customers who were exposed to ads are **incremental**. That is, none of the sales would have happened without the advertising.
- The duration of advertising response is relatively short and the ad exposures are well spaced out over time, so that there are not other channels providing an “assist”.

These conditions might hold for a startup or a new product, but not for too many other advertisers.



Source: [Mazeberry blog](#)

Last touch is largely dead

There are many vendors such as Convertro, VisualIQ and MarketShare that have provided model-based attribution solutions.



Source: [Google Analytics 360](#)

In Google Analytics 360, released in 2016, Google also abandoned rule-based attribution in favor of model-based approaches.

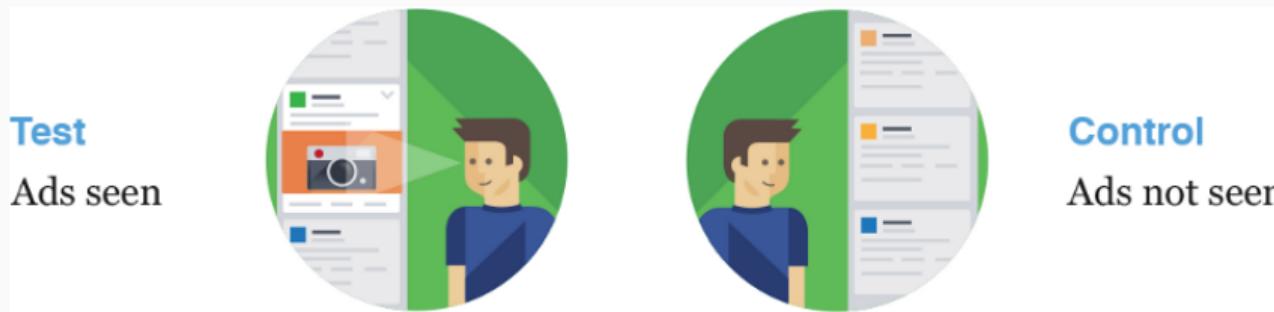
Holdout testing

Holdout testing

What is a holdout test?

Holdout testing

The gold standard for measuring incremental sales is an experiment, where we **randomly** assign customers to be exposed or not exposed to an ad.



Source: [Designing with Science, medium.com](https://designingwithscience.com)

For example, with email we can take our list of target emails and randomly select a set of them to not receive an email. Facebook also has a tool for randomized holdouts (for large advertisers.)

Why holdout testing works

The magic behind a holdout test is the randomization. Another name for holdout testing is **randomized controlled trial**.

Why holdout testing works

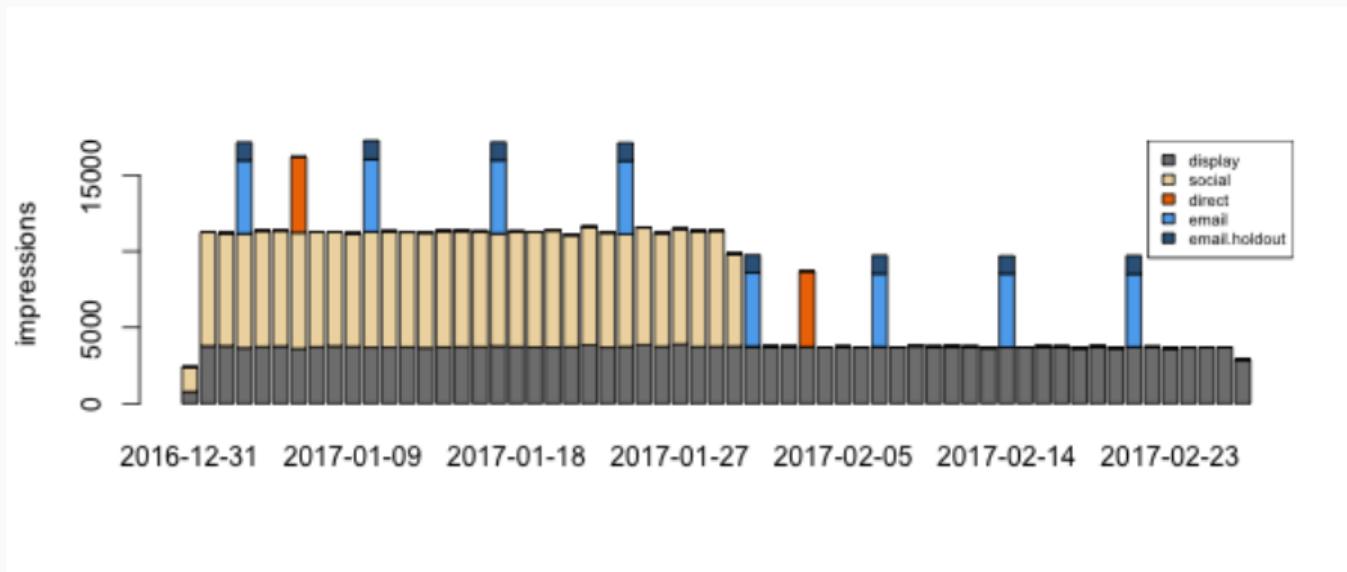
The magic behind a holdout test is the randomization. Another name for holdout testing is **randomized controlled trial**.

By randomly selecting the customers for the control group, we assure that the two groups are the same on average. The treatment and control groups will be similar in their propensity to purchase and response to ads. Statisticians call this **probabilistically equivalent**.

Holdout testing

Holdout test analysis in R

Email holdout tests in the example data



There was an email on 2017-01-03 that included a holdout group. Let's analyze this test.

Function for analyzing holdout tests

```
holdout.test <- function(test.date, delay=0, window, impress, trans) {  
  test.ids <- unique(impress$id[impress$channel=="email" & impress$date==test.date])  
  control.ids <- unique(impress$id[impress$channel=="email.holdout" & impress$date==test.date])  
  tdata <- data.frame(ids=c(test.ids, control.ids))  
  tdata$group[tdata$id %in% test.ids] <- "test"  
  tdata$group[tdata$id %in% control.ids] <- "control"  
  in.window <- trans$date>=(test.date+delay) & trans$date<(test.date+window+delay)  
  tdata$convert <- tdata$id %in% trans$id[in.window]  
  ttable <- xtabs(~ group + convert, data=tdata)  
  mosaicplot(~ group + convert, data=tdata,  
             main=paste("Email Test on", test.date))  
  proptest <- prop.test(x=ttable[, "TRUE"], n=xtabs(~ group, data=tdata))  
  diff.conv <- c(diff=proptest$estimate[2]-proptest$estimate[1], ci=-proptest$conf.int)  
  out <- list(diff.conv, ttable, proptest)  
}
```

Analysis of email holdout test

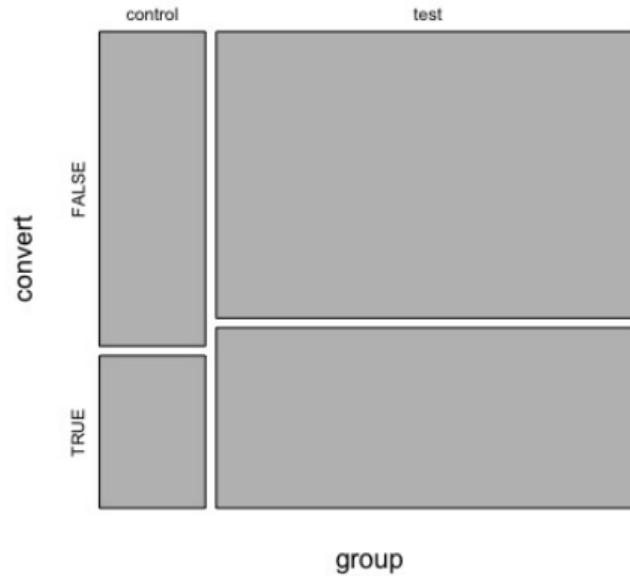
```
> (holdout.test(test.date=as.Date("2017-01-03"), window=7, impress=impress,
+                trans=trans))
[[2]]
    convert
group      FALSE  TRUE
control     810   393
test        2944  1854

[[3]]
2-sample test for equality of proportions with
continuity correction

data: ttable[, "TRUE"] out of xtabs(~group, data = tdata)
X-squared = 14.395, df = 1, p-value = 0.0001482
alternative hypothesis: two.sided
95 percent confidence interval:
-0.09011755 -0.02933788
sample estimates:
prop 1     prop 2
0.3266833 0.3864110
```

Reporting of email holdout test

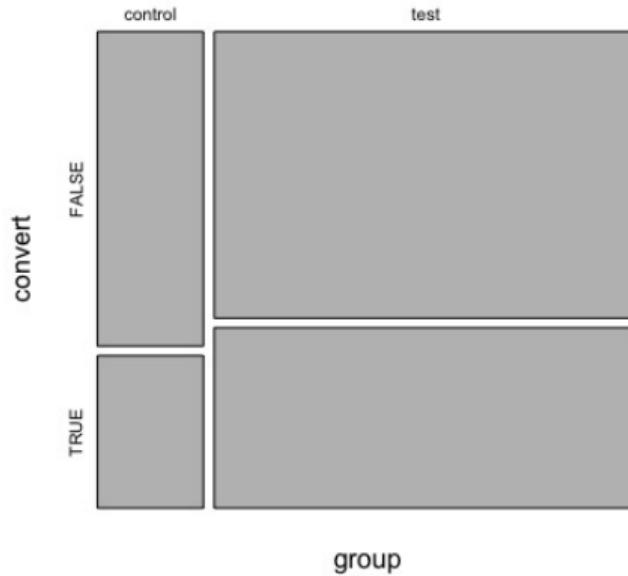
Email Test on 2017-01-03



The test group had a 38.6% conversion rate in the 7 days after the email was sent, versus a 32.7% conversion rate for the control group.

Reporting of email holdout test

Email Test on 2017-01-03



The test group had a 38.6% conversion rate in the 7 days after the email was sent, versus a 32.7% conversion rate for the control group.

The email on 2017-01-03 produced incremental sales. The incremental increase in conversion rate is between +2.9% and +9.0% (95% CI).

Another analysis of email holdout test on 2017-01-01

```
> (holdout.test(test.date=as.Date("2017-01-03"), window=3, impress=impress, trans=trans))
[[2]]
      convert
group      FALSE  TRUE
control    1063   140
test       3850   948

[[3]]
 2-sample test for equality of proportions with continuity correction

data: ttable[, "TRUE"] out of xtabs(~group, data = tdata)
X-squared = 42.187, df = 1, p-value = 0.0000000008295
alternative hypothesis: two.sided
95 percent confidence interval:
 -0.10306428 -0.05934891
sample estimates:
 prop 1     prop 2 
0.1163757 0.1975823
```

Setting the response window

If you change the response window, you will get a different answer about how much advertising increases sales.

Setting the response window

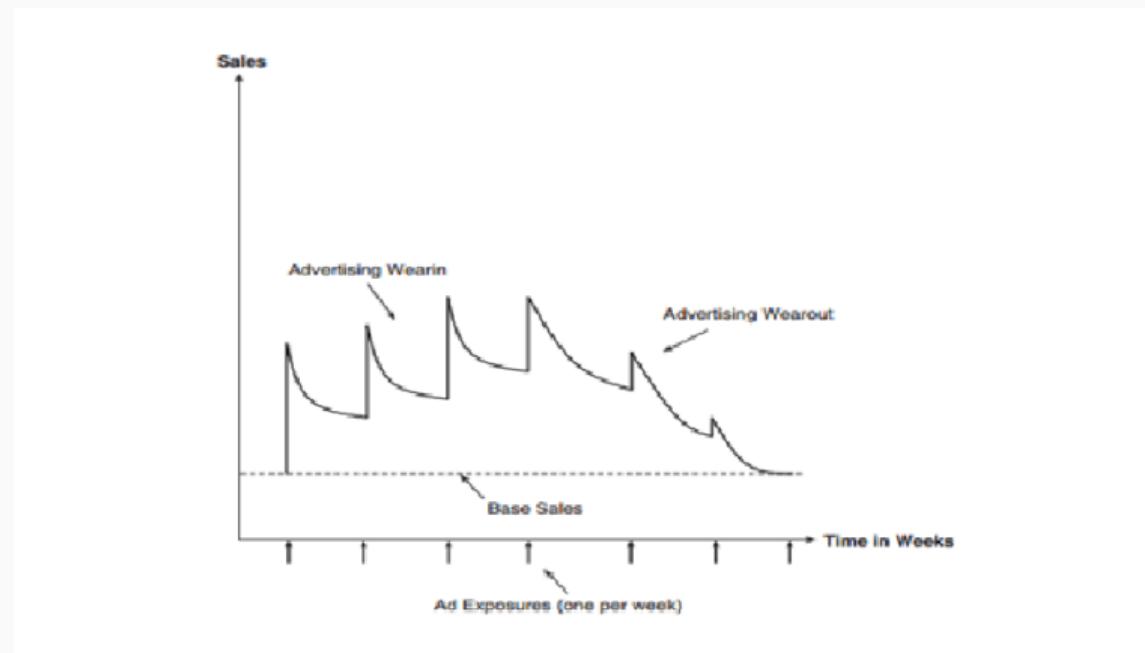
If you change the response window, you will get a different answer about how much advertising increases sales.

- For the email test on 2017-01-03, a response window of 3 days indicated a increase in conversion of 8.1% (95% CI = [5.9, 10.3]) versus 6.0% (95% CI = 2.9, 9.0) for a 7 day window.

What is happening here?

Checking ad response over time

Ad response is often greatest just after exposure and then falls off over time.



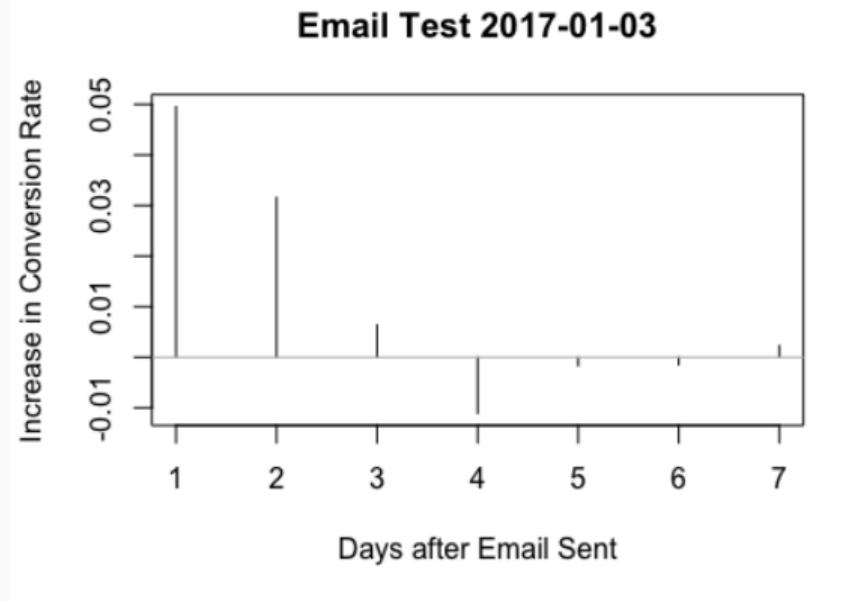
Source: Michael Wolfe

Ad response over time for 2017-01-03 email test

With a holdout test, you can study the ad response over time to learn what the pattern looks like over time.

```
day1 <- holdout.test(test.date=as.Date("2017-01-03"), delay=0, window=1,
impress=impress, trans=trans)
...
day7 <- holdout.test(test.date=as.Date("2017-01-03"), delay=6, window=1,
impress=impress, trans=trans)
incr.conv <- c(day1[[1]][1], day2[[1]][1], day3[[1]][1], day4[[1]][1], day5[[1]][1],
               day6[[1]][1], day7[[1]][1])
plot(incr.conv, type="h", xlab="Days after Email Sent", ylab="Increase in Conversion Rate")
abline(h=0)
```

Ad response over time for 2017-01-03 email test



In the test on 2017-01-03, the lift in conversion rate falls to zero in about three days.

Other email tests

For those following along in R, try analyzing any of the other tests by modifying the code.

Dates of tests are:

2017-01-17

2017-01-24

2017-01-31

2017-02-07

2017-02-14

2017-02-21

Holdout testing

Planning a holdout test

Designing a holdout test

How do you measure response? For how long?

- Conversion (binary), Sales (number), or Site Visits (count)
- During or after campaign

What advertising do the treatment and control groups receive?

- How many ads does the treatment group see? Usual? Heavy up?
- What is used for the control? “Go dark”? alternative ad? Usual?

Which customers are included in the test?

- Current customers, good customers or prospects?

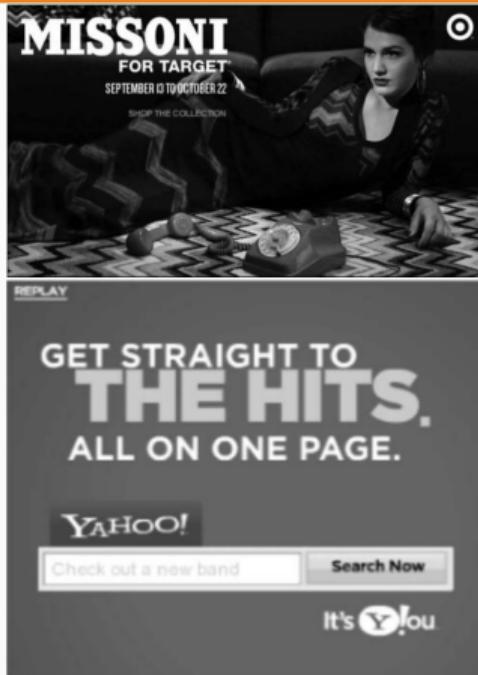
How should customers be assigned to treatment and control groups?

- Always random, but doesn't have to be an equal split.

Holdout testing

**Example: Johnson, Lewis and Reiley
(2017)**

Display advertising holdout test at Yahoo!



During two weeks in Spring 2010, a national apparel retailer ran a test on Yahoo! to measure response to **display ads** on Yahoo!.

The ads displayed the retailer brand and brands of apparel firm, with slideshow transitions between photographs and text.

This test is described in detail in [Johnson, Lewis and Reiley \(2017\)](#) "When less is more: Data and power in advertising experiments", *Marketing Science* 36(1), 43-53.

Source: [Johnson, Lewis and Reiley \(2017\)](#)

Display advertising holdout test: design

How do you measure response? For how long?

- Sales (\$) at the target retailer during the two week campaign

Display advertising holdout test: design

How do you measure response? For how long?

- Sales (\$) at the target retailer during the two week campaign

What advertising do the treatment and control groups receive?

- **Full**: all retailer ads at all opportunities
- **Control**: house Yahoo! ads at all opportunities
- **Half**: half retailer and half house ads

Display advertising holdout test: design

How do you measure response? For how long?

- Sales (\$) at the target retailer during the two week campaign

What advertising do the treatment and control groups receive?

- **Full**: all retailer ads at all opportunities
- **Control**: house Yahoo! ads at all opportunities
- **Half**: half retailer and half house ads

Which customers are included in the test?

- Registered Yahoo! users who were also customers of the retailer were assigned to groups at the start of the two week period. Many of these users never visit Yahoo! during the campaign and don't have the opportunity to see the ads.
- Database match allows us to track users between Yahoo! and the retailer.

Display advertising holdout test: design

How do you measure response? For how long?

- Sales (\$) at the target retailer during the two week campaign

What advertising do the treatment and control groups receive?

- **Full**: all retailer ads at all opportunities
- **Control**: house Yahoo! ads at all opportunities
- **Half**: half retailer and half house ads

Which customers are included in the test?

- Registered Yahoo! users who were also customers of the retailer were assigned to groups at the start of the two week period. Many of these users never visit Yahoo! during the campaign and don't have the opportunity to see the ads.
- Database match allows us to track users between Yahoo! and the retailer.

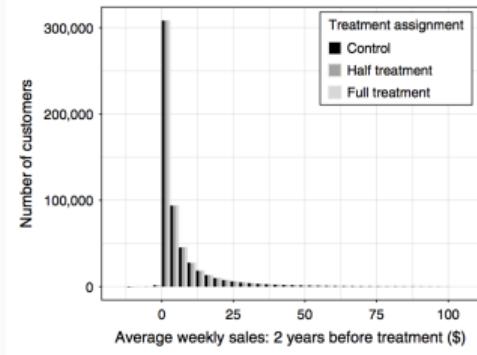
How should customers be assigned to treatment and control groups?

- Randomly, with one third in each group

Display advertising holdout test: randomization checks

It is a good idea to check that the treatment and control groups looked similar prior to the test.

	Treatment group			<i>p</i> -value
	Full	Half	Control	
Sample size	1,032,204	1,032,074	1,032,299	0.988
Female (mean)	68.5%	68.5%	68.5%	0.794
Age (mean)	43.6	43.6	43.6	0.607
Yahoo! page views ^a (mean)	245.8	244.4	243.5	0.132 ^d
Pretreatment sales (two years, mean)	\$857.74	\$859.30	\$855.54	0.475
Pretreatment sales (two weeks, mean)	\$19.34	\$19.24	\$19.10	0.517
Treated subsample				
Exposed sample	572,574	571,222	570,908	0.254
Yahoo! page views (mean)	412.2	411.5	410.1	0.108
Ad views (mean)	33.42	33.41	33.66	0.164
Ad views (median)	15	15	15	
Retailer ad views (mean)	33.42	16.69	—	0.801
Control ad views (mean)	—	16.72	33.66	0.165
Retailer ad click-through rate ^b (%)	0.19	0.24	—	
Retailer ad clicker rate ^c (%)	4.91	3.39	—	



Source: Johnson, Lewis and Reiley (2017)

The three groups appear to be similar. So, the randomization worked and we are comparing “apples to apples”.

Display advertising holdout test: result

Sales Increase for Full and Half Treatments (\$)

	(1)	(2)
Subset of users ^a	Everyone	Treated
Sales after first ad exposure ^b		
Full treatment (\$)	0.673** (0.317)	0.525** (0.237)
Half treatment (\$)	0.0248 (0.311)	0.189 (0.235)
Constant (\$)	15.52*** (0.122)	15.53*** (0.166)

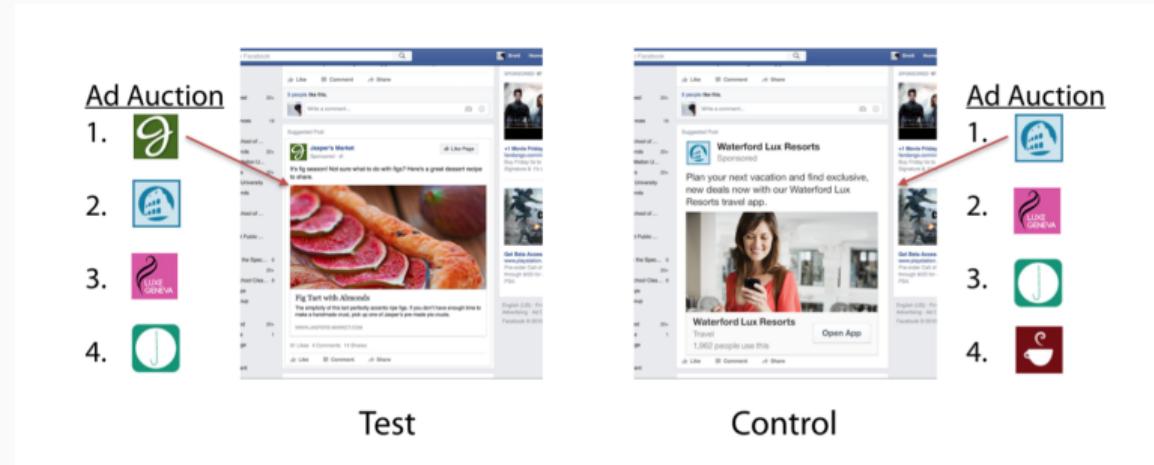
Looking at all the customers, the treatment increased sales by \$0.67 (95% CI = [0.05, 1.29]). Among those who actually were exposed to ads, the effect was \$0.52 (95% CI = [0.06, 0.97]).

Holdout testing

**Example: Gordon, Zettelmeyer,
Bahargava and Chapsky (2016WP)**

Social advertising holdout test

During two weeks in the first half of 2015, an omni-channel retailer conducted a test to determine the lift Facebook ads in the newsfeed.



Source: Gordon et al. (2016 WP)

This test is reported in Gordon, Zettelmeyer, Bhargava and Chapsky (2016WP) "A Comparison of Approaches to Advertising Measurement: Evidence from Big Field Experiments at Facebook"

Search advertising holdout test design

Which customers are included in the test?

- Facebook users that met the targeting criteria for the campaign

Search advertising holdout test design

Which customers are included in the test?

- Facebook users that met the targeting criteria for the campaign

How do you measure response? For how long?

- Purchases at the retailers website measured via Facebook conversion pixel during the campaign period and up to several weeks after the study ended.

Search advertising holdout test design

Which customers are included in the test?

- Facebook users that met the targeting criteria for the campaign

How do you measure response? For how long?

- Purchases at the retailers website measured via Facebook conversion pixel during the campaign period and up to several weeks after the study ended.

How should customers be assigned to treatment and control groups?

- Randomly, 70% test and 30% control

Search advertising holdout test design

Which customers are included in the test?

- Facebook users that met the targeting criteria for the campaign

How do you measure response? For how long?

- Purchases at the retailers website measured via Facebook conversion pixel during the campaign period and up to several weeks after the study ended.

How should customers be assigned to treatment and control groups?

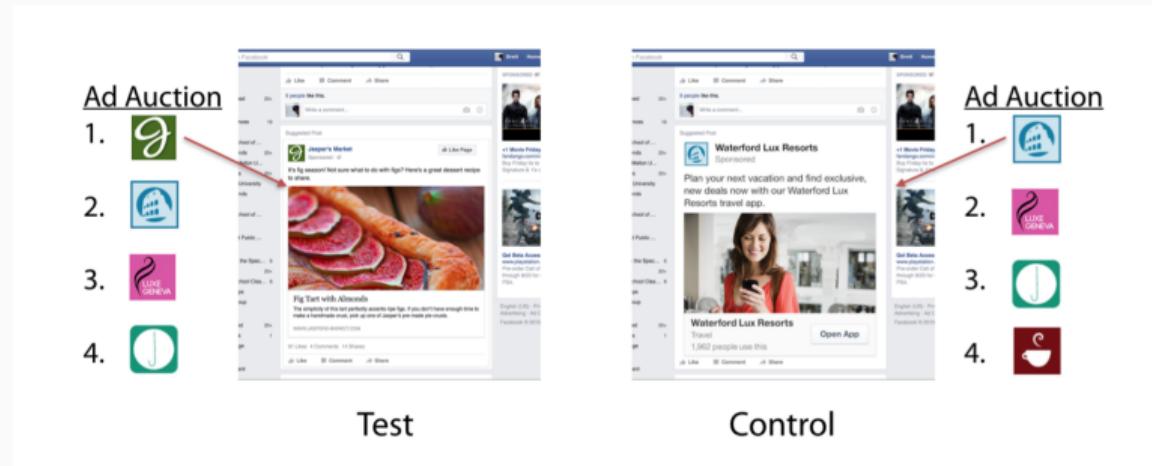
- Randomly, 70% test and 30% control

What advertising do the treatment and control groups receive?

- Test: user is eligible to see ads and will if the retailer wins the Facebook auction
- Control: competitor ads

Social advertising holdout test: Control condition

Since purchasing house ads or public service announcements (PSAs) is expensive, the Facebook experiment ran a competitor ad for the control.



Source: Gordon et al. (2016 WP)

By using the “next advertiser” as the control condition, the sales lift we estimate is relative to “what would have happened had we not done the campaign.”

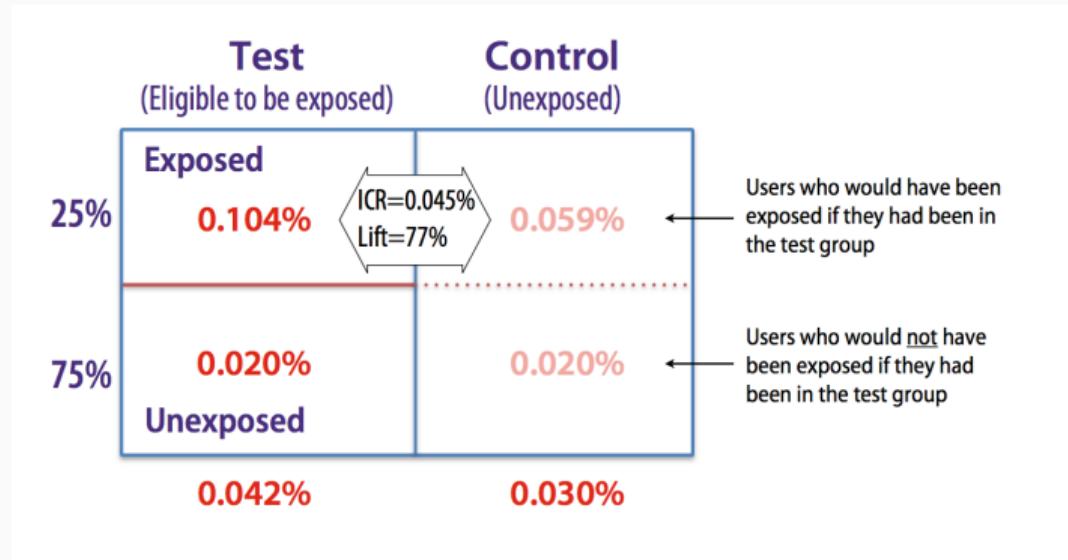
Social holdout test: Randomization check

Variable	Control group	Test group	p-value
Average user age	31.7	31.7	0.33
% of users who are male	17.2%	17.2%	0.705
Length of time using FB (days)	2,288	2,287	0.24
% of users with status “married”	19.6	19.6	0.508
% of users status “engaged”	13.8	13.8	0.0892
% of users status “single”	14.0	14.0	0.888
# of FB friends	485.7	485.7	0.985
# of FB uses in last 7 days	6.377	6.376	0.14
# of FB uses in last 28 days	25.5	25.5	0.172

Source: Gordon et al. (2016 WP)

The test and control groups appear to be similar. Randomization worked.

Social holdout test: Results



Source: Gordon et al. (2016 WP)

There is a 0.045% increase in conversions when comparing the users exposed to ads to the users in the control group that would have been exposed.

Facebook holdout testing tool: Conversion Lift

What is conversion lift?

Conversion lift accurately captures the impact that Facebook ads have in driving business for marketers. Here's how it works:

1. When creating a Facebook campaign, a randomized test group (people that see ads) and control group (people that don't) are established.
2. The advertiser securely shares conversion data from the campaign with Facebook. Typically, this data comes from sources like the Facebook Custom Audiences pixel, conversion pixel or secure point-of-sale (POS) data.
3. Facebook determines additional lift generated from the campaign by comparing conversions in the test and control groups.
4. The results of the study are made available in Ads Manager.

Source: [Facebook Conversion Lift](#)

There were **5,000** additional buyers as a result of your ads, an increase of **3.00%**.



+ 5,000
Incremental Buyers



+ 3.00%
Conversion Lift



GROUP SIZE

CONVERSION RATE

BUYERS

TEST

800K



25.0%



200K

CONTROL

200K

4.0x

SCALED
CONTROL

800K



24.4%



195K

INCREMENTAL
BUYERS

5,000 ± 1,102

STEP 1: COMPARE THE GROUPS

Scaling the control group to the same size as the test group provides an equal comparison between groups.

STEP 2: APPLY THE CONVERSION RATES

Applying the conversion rate from the control group to the scaled control shows how many people would have bought if the control group were the same size as the test group.

STEP 3: FIND THE INCREMENTAL BUYERS

Subtracting the buyers in the scaled control group from the buyers in the test group gives you the incremental buyers, the number of buyers directly attributable to your ads.

[View Conversion Lift >](#)

[Close this card](#)

Holdout testing

**Example: Blake, Nosko and Tadellis
(2015)**

Search advertising holdout test

The screenshot shows a Google search results page for the query "used gibson les paul". The top result is a sponsored ad from "Used Guitar - Used Gear in Like New Condition." Below it is another sponsored ad for "Used Les Paul Gibson". Organic search results include links to "Gibson Les Paul Used on eBay", "Shop for used gibson les paul on Google", "Gibson | Dave's Guitar Shop", "Gibson Guitar - Get great deals for Gibson Guitar on eBay!", and "Gibson Les Paul - eBay - Find Popular Products on eBay!". Each result includes a thumbnail image of a guitar and its price.

In April-July 2012, eBay ran a holdout experiment to determine the sales lift of their search advertising targeting keywords other than “eBay”.

This test is reported in detail in [Blake, Nosko and Tadelis \(2015\) “Consumer Heterogeneity and Paid Search Effectiveness: A Large-Scale Field Experiment”, *Econometrica* 83\(1\), 155-174.](#)

Source: [Blake, Nosko and Tadelis \(2015\)](#)

Search advertising holdout test: Design

Which customers are included in the test?

- Google AdWords does not provide a testing tool that will randomize search ad exposure at the user level. Instead, they conducted a **geo-test** with 30% of the DMAs in the US market.

Search advertising holdout test: Design

Which customers are included in the test?

- Google AdWords does not provide a testing tool that will randomize search ad exposure at the user level. Instead, they conducted a **geo-test** with 30% of the DMAs in the US market.

How do you measure response? For how long?

- eBay sales (\$) shipped to the DMA

Search advertising holdout test: Design

Which customers are included in the test?

- Google AdWords does not provide a testing tool that will randomize search ad exposure at the user level. Instead, they conducted a **geo-test** with 30% of the DMAs in the US market.

How do you measure response? For how long?

- eBay sales (\$) shipped to the DMA

What advertising do the treatment and control groups receive?

- Test: keep search advertising at usual level
- Control: “go dark”

Search advertising holdout test: Design

Which customers are included in the test?

- Google AdWords does not provide a testing tool that will randomize search ad exposure at the user level. Instead, they conducted a **geo-test** with 30% of the DMAs in the US market.

How do you measure response? For how long?

- eBay sales (\$) shipped to the DMA

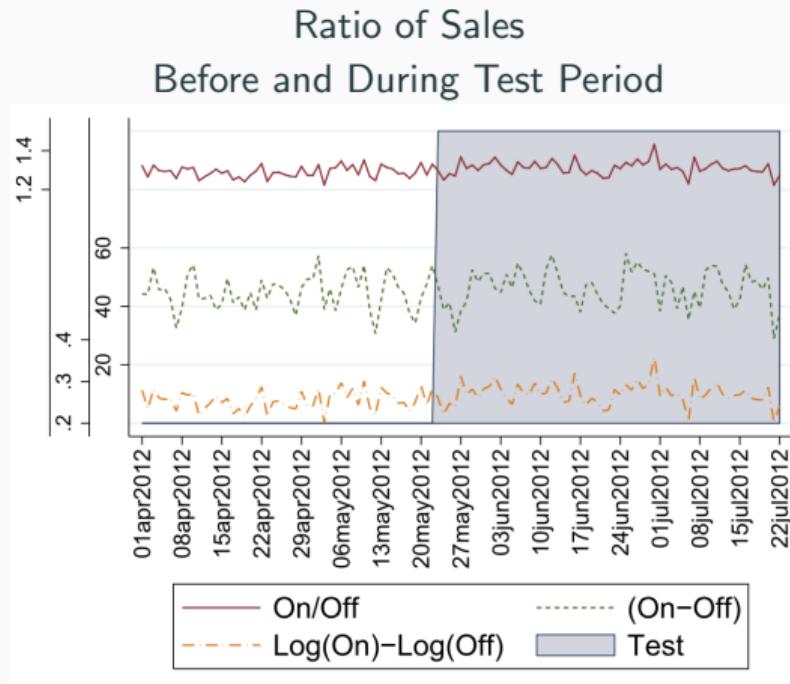
What advertising do the treatment and control groups receive?

- Test: keep search advertising at usual level
- Control: “go dark”

How should DMAs be assigned to treatment and control groups?

- DMAs were **matched** based on prior eBay sales and then randomly assigned to groups.

Search advertising holdout test: Findings

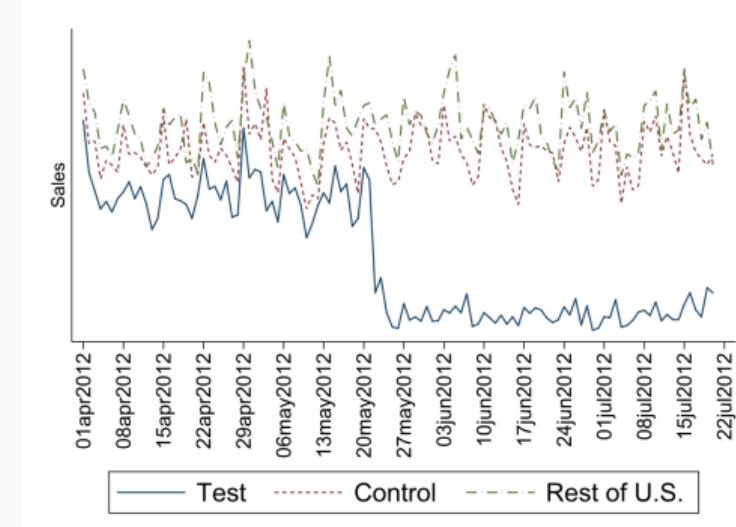


Search advertising does not increase eBay sales.
As a result, eBay has drastically reduced search advertising.

Source: [Blake, Nosko and Tadellis \(2015\)](#)

Search advertising holdout test: “Last Touch”

“Last Touch” Attributed Sales
(sale within 24 hours of search ad exposure)
Before and During Test Period



If you were using last touch attribution, you would think that search advertising made a big difference.

Source: [Blake, Nosko and Tadellis \(2015\)](#)

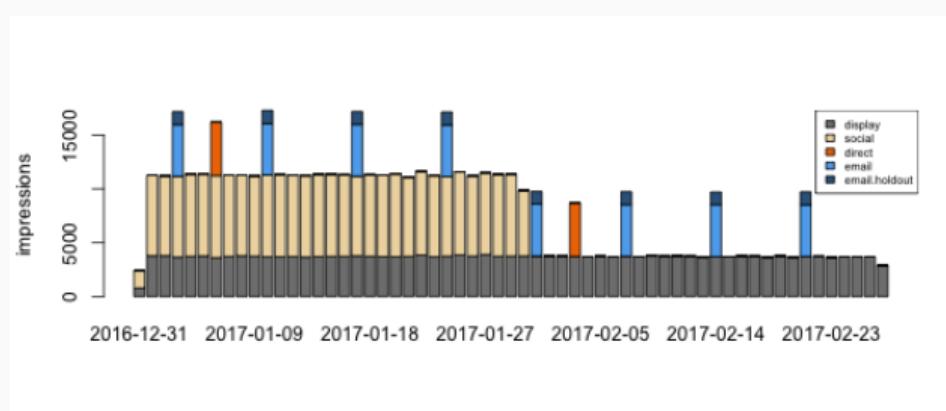
Marketing mix modeling

Marketing mix modeling

What is marketing mix modeling?

Marketing mix modeling

Before advertisers had access to user-level ad exposure data (like our example data set) they only had data on what they had spent on advertising and (sometimes) the number of people who saw the ad. This type of spending data is similar to the aggregate impressions data we looked at for our data set.



To find the correlation between **total sales** in each day/week/month to **advertising spending or impressions** on that same day/week/month, you estimate **regression** model.

A simple marketing mix model

A regression is an equation relating a response ("dependent variable") to one or more other variables ("independent variables").

A simple marketing mix model

A regression is an equation relating a response ("dependent variable") to one or more other variables ("independent variables").

A simple marketing mix model is:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

A simple marketing mix model

A regression is an equation relating a response ("dependent variable") to one or more other variables ("independent variables").

A simple marketing mix model is:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

The words represent data we have for each time period and the β 's represent the unknown relationship between ad impressions and sales. For example β_1 is the increase in sales we get for each additional display impression.

A simple marketing mix model

A regression is an equation relating a response ("dependent variable") to one or more other variables ("independent variables").

A simple marketing mix model is:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

The words represent data we have for each time period and the β 's represent the unknown relationship between ad impressions and sales. For example β_1 is the increase in sales we get for each additional display impression.

Using observations of the sales and advertising data, when can estimate these unknown parameters of the model.

Marketing mix modeling

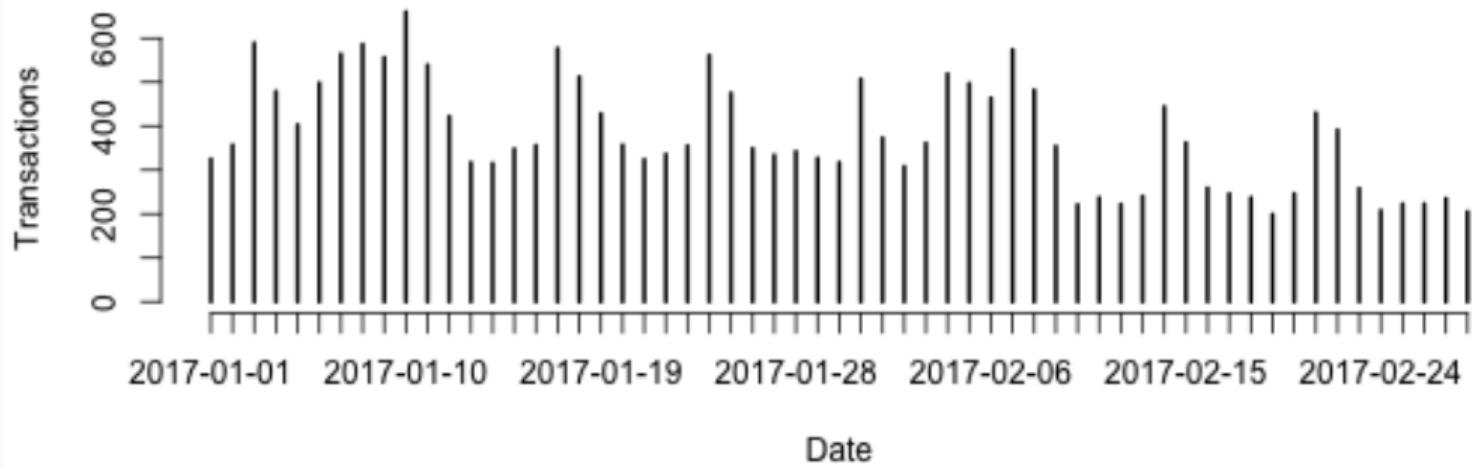
Marketing-mix modeling in R

Aggregating the user-level sales data

As a reminder, we already computed the aggregate transactions for each day from our raw user-level data using a crosstab.

```
> (transbyday <- xtabs(~date, data=trans))
date
2017-01-01 2017-01-02 2017-01-03 2017-01-04 2017-01-05 2017-01-06 2017-01-07
      325        357        589        479        403        498        564
2017-01-08 2017-01-09 2017-01-10 2017-01-11 2017-01-12 2017-01-13 2017-01-14
      586        556        660        539        422        317        315
2017-01-15 2017-01-16 2017-01-17 2017-01-18 2017-01-19 2017-01-20 2017-01-21
      348        356        577        512        428        357        324
```

Sales by day



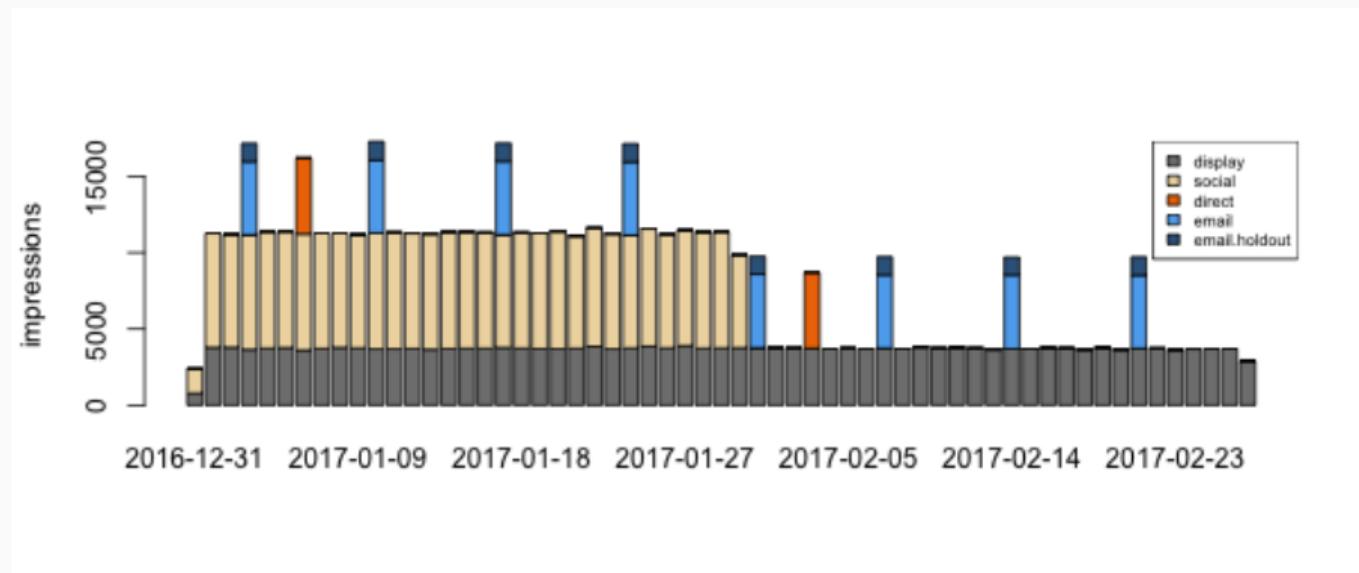
Aggregating the user-level advertising data

We also computed the aggregate number of impressions on each day.

```
> (cadence <- xtabs(~date+channel, data=impress))
```

	channel	direct	display	email	email.holdout	social
date						
2016-12-31		0	788	0	0	1610
2017-01-01		0	3786	0	0	7481
2017-01-02		0	3792	0	0	7416
2017-01-03		0	3656	4798	1203	7505
2017-01-04		0	3731	0	0	7648
2017-01-05		0	3770	0	0	7620
2017-01-06	4974	3611	0	0	0	7614
2017-01-07		0	3719	0	0	7552
2017-01-08		0	3780	0	0	7504
2017-01-09		0	3744	0	0	7446

Advertising by day



Combining the sales and impressions data

When we fit models, it will be convenient to put the sales and impressions data together in the same data frame.

```
> mdata <- as.data.frame(cbind(trans=transbyday[1:57], cadence[2:58,])) # aligning  
> head(mdata)  
      trans direct display email email.holdout social  
2017-01-01    325      0    3786      0            0  7481  
2017-01-02    357      0    3792      0            0  7416  
2017-01-03    589      0    3656    4798          1203  7505  
2017-01-04    479      0    3731      0            0  7648  
2017-01-05    403      0    3770      0            0  7620  
2017-01-06    498   4974    3611      0            0  7614
```

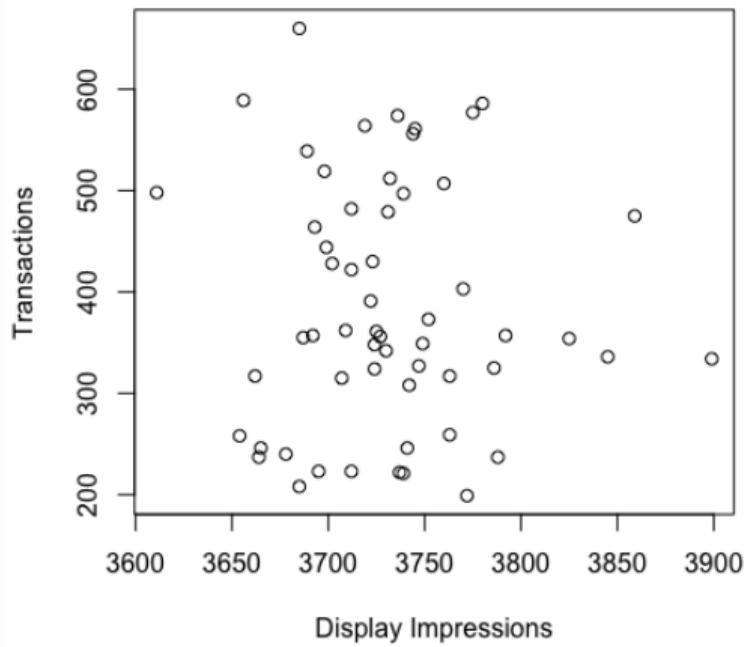
Data exploration

Now that we have the data aggregated, we should do some exploration before modeling.

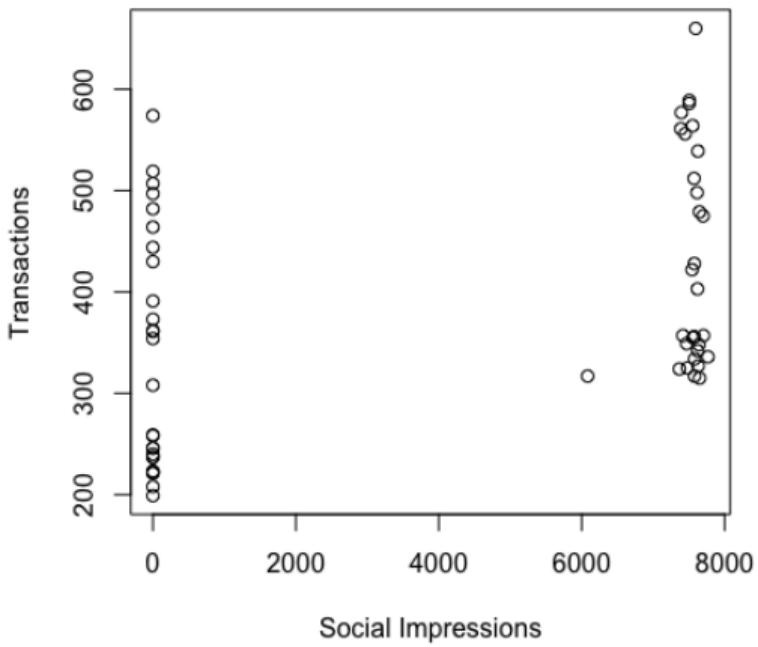
```
> plot(x=mdata$direct, y=mdata$trans, xlab="Direct Impressions",
+       ylab="Transactions", main="Direct Impressions v. Transactions")
> plot(x=mdata$email, y=mdata$trans, xlab="Email Impressions",
+       ylab="Transactions", main="Email Impressions v. Transactions")
> plot(x=mdata$display, y=mdata$trans, xlab="Display Impressions",
+       ylab="Transactions", main="Display Impressions v. Transactions")
> plot(x=mdata$social, y=mdata$trans, xlab="Social Impressions",
+       ylab="Transactions", main="Social Impressions v. Transactions")
```

Daily display and social impressions versus transactions

Display Impressions v. Transactions

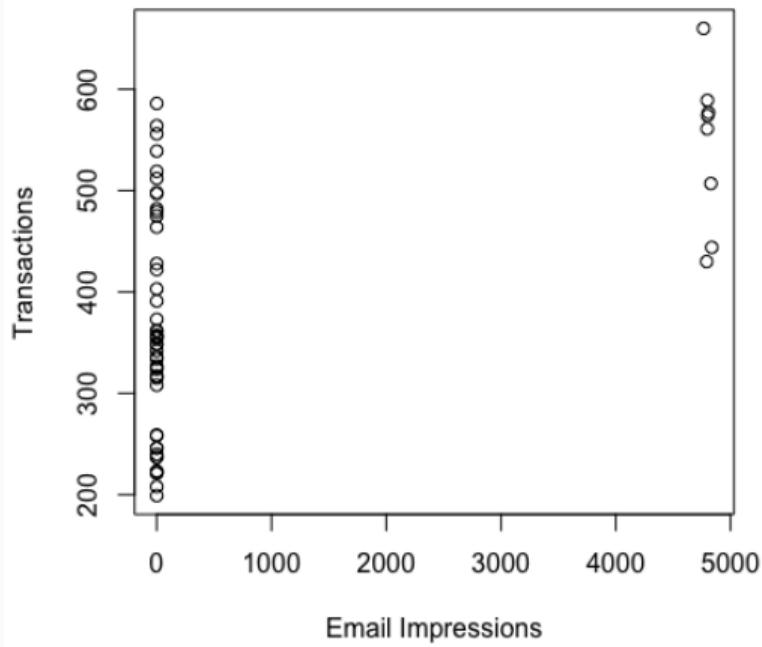


Social Impressions v. Transactions

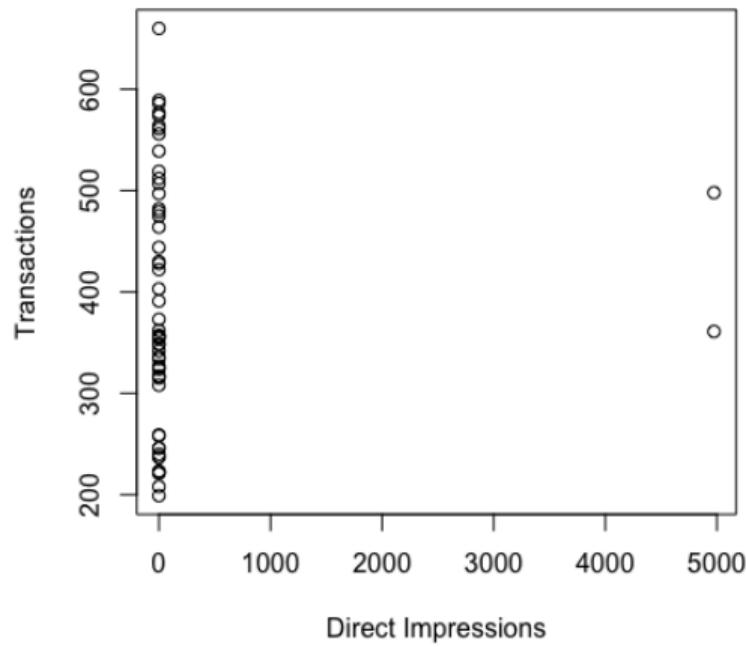


Daily email and direct impressions versus transactions

Email Impressions v. Transactions



Direct Impressions v. Transactions



Running a simple media mix model

```
> m1 <- lm(trans~direct+social+email+display+social+email.holdout, data=mdata)
> summary(m1)
...
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 438.501429  940.622551   0.466 0.643070
direct       0.014529    0.013853   1.049 0.299206
social       0.012663    0.003379   3.747 0.000457 ***
email        -0.161643   0.321735  -0.502 0.617541
display      -0.035299   0.252371  -0.140 0.889313
email.holdout  0.805475   1.290198   0.624 0.535211
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 92.53 on 51 degrees of freedom
Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4119
F-statistic: 8.845 on 5 and 51 DF,  p-value: 0.00000427
```

Things to notice in the model output

```
> m1 <- lm(trans~direct+social+email+display+social+email.holdout,  
           data=mdata)  
> summary(m1)  
  
Call:  
lm(formula = trans ~ direct + social + email + display + social +  
    email.holdout, data = mdata)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-106.35   -67.79   -41.56    54.42   211.04  
  
Coefficients:  
            Estimate Std. Error t value Pr(>|t|)  
(Intercept) 438.501429  940.622551   0.466  0.643070  
direct       0.014529   0.013853   1.049  0.299206  
social       0.012663   0.003379   3.747  0.000457 ***  
email        -0.161643   0.321735  -0.502  0.617541  
display      -0.035299   0.252371  -0.140  0.889313  
email.holdout 0.805475   1.290198   0.624  0.535211  
---  
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1  
  
Residual standard error: 92.53 on 51 degrees of freedom  
Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4119  
F-statistic: 8.845 on 5 and 51 DF,  p-value: 0.00000427
```

The only statistically significant effect is for social impressions, and we get 0.012663 additional transactions for each social impression.

Things to notice in the model output

```
> m1 <- lm(trans~direct+social+email+display+social+email.holdout,  
           data=mdata)  
  
> summary(m1)  
  
Call:  
lm(formula = trans ~ direct + social + email + display + social +  
   email.holdout, data = mdata)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-106.35 -67.79 -41.56  54.42 211.04  
  
Coefficients:  
            Estimate Std. Error t value Pr(>|t|)  
(Intercept) 438.501429  940.622551   0.466  0.643070  
direct       0.014529   0.013853   1.049  0.299206  
social       0.012663   0.003379   3.747  0.000457 ***  
email        -0.161643   0.321735  -0.502  0.617541  
display      -0.035299   0.252371  -0.140  0.889313  
email.holdout 0.805475   1.290198   0.624  0.535211  
---  
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1  
  
Residual standard error: 92.53 on 51 degrees of freedom  
Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4119  
F-statistic: 8.845 on 5 and 51 DF,  p-value: 0.00000427
```

The only statistically significant effect is for social impressions, and we get 0.012663 additional transactions for each social impression.

The standard error for display is very large, which means we don't have a precise estimate of the effect of display. This happened because daily display impressions are pretty much the same every day. The data is not informative!

Things to notice in the model output

```
> m1 <- lm(trans~direct+social+email+display+social+email.holdout,
   data=mdata)
> summary(m1)

Call:
lm(formula = trans ~ direct + social + email + display + social +
    email.holdout, data = mdata)

Residuals:
    Min      1Q  Median      3Q     Max 
-106.35 -67.79 -41.56  54.42 211.04 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 438.501429  940.622551   0.466  0.643070    
direct       0.014529   0.013853   1.049  0.299206    
social       0.012663   0.003379   3.747  0.000457 ***  
email        -0.161643   0.321735  -0.502  0.617541    
display      -0.035299   0.252371  -0.140  0.889313    
email.holdout 0.805475   1.290198   0.624  0.535211    
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 92.53 on 51 degrees of freedom
Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4119 
F-statistic: 8.845 on 5 and 51 DF,  p-value: 0.00000427
```

The only statistically significant effect is for social impressions, and we get 0.012663 additional transactions for each social impression.

The standard error for display is very large, which means we don't have a precise estimate of the effect of display. This happened because daily display impressions are pretty much the same every day. The data is not informative!

The estimated effect of email and display is negative (but not significant).

Adding other variables

In addition to advertising impressions, we might want to include other predictors of transactions in our model, such as the day of the week.

```
> mdata$dayofweek <- weekdays(as.Date(rownames(mdata)))
> m2 <- lm(trans ~ email + direct + display + social + dayofweek, data = mdata)
```

Our second model

```
> summary(m2)

Call:
lm(formula = trans ~ email + direct + display + social + dayofweek,
    data = mdata)

Residuals:
    Min      1Q  Median      3Q     Max 
-87.09 -51.78 -21.70  25.46 221.46 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 348.192012  891.949059   0.390 0.698063  
email        -0.960694   1.463482  -0.656 0.514809  
direct        0.029537   0.014165   2.085 0.042632 *  
display       -0.031124   0.239069  -0.130 0.896986  
social         0.012754   0.003077   4.144 0.000145 *** 
dayofweekMonday 71.405198  45.616341   1.565 0.124357  
dayofweekSaturday 64.445464  45.601326   1.413 0.164318  
dayofweekSunday  55.293919  44.598357   1.240 0.221330  
dayofweekThursday 67.663736  45.341037   1.492 0.142440  
dayofweekTuesday 4877.229334 7029.997193  0.694 0.491313  
dayofweekWednesday 171.086874  45.336050   3.774 0.000459 *** 
...
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 83.94 on 46 degrees of freedom
Multiple R-squared:  0.6024,    Adjusted R-squared:  0.516 
F-statistic:  6.97 on 10 and 46 DF,  p-value: 0.000001582
```

When we add other variables, it can change the coefficients of the model.

Our second model

```
> summary(m2)

Call:
lm(formula = trans ~ email + direct + display + social + dayofweek,
    data = mdata)

Residuals:
    Min      1Q  Median      3Q     Max 
-87.09 -51.78 -21.70  25.46 221.46 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 348.192012  891.949059  0.390 0.698063  
email        -0.960694   1.463482  -0.656 0.514809 *  
direct        0.029537   0.014165  2.085 0.042632 *  
display       -0.031124   0.239069  -0.130 0.896986  
social         0.012754   0.003077  4.144 0.000145 *** 
dayofweekMonday 71.405198  45.616341  1.565 0.124357  
dayofweekSaturday 64.445464  45.601326  1.413 0.164318  
dayofweekSunday  55.293919  44.598357  1.240 0.221330  
dayofweekThursday 67.663736  45.341037  1.492 0.142440  
dayofweekTuesday 4877.229334 7029.997193  0.694 0.491313  
dayofweekWednesday 171.086874  45.336050  3.774 0.000459 *** 
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 83.94 on 46 degrees of freedom
Multiple R-squared:  0.6024,    Adjusted R-squared:  0.516 
F-statistic:  6.97 on 10 and 46 DF,  p-value: 0.000001582
```

When we add other variables, it can change the coefficients of the model.

With our new model we find a significant association between direct impressions and transactions (29.5 additional transactions / 1000 direct impressions) in addition to social (12.8 additional transactions / 1000 impressions).

What is our model missing?

Our story about where our data comes from says:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

What is our model missing?

Our story about where our data comes from says:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

Our model assumes that an impression on day t has an effect on the number of transactions on day t and that those impressions have no effect on sales on other days.

What is our model missing?

Our story about where our data comes from says:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

Our model assumes that an impression on day t has an effect on the number of transactions on day t and that those impressions have no effect on sales on other days.

Do we really believe that ads only work on the day they were served/sent?

What is our model missing?

Our story about where our data comes from says:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

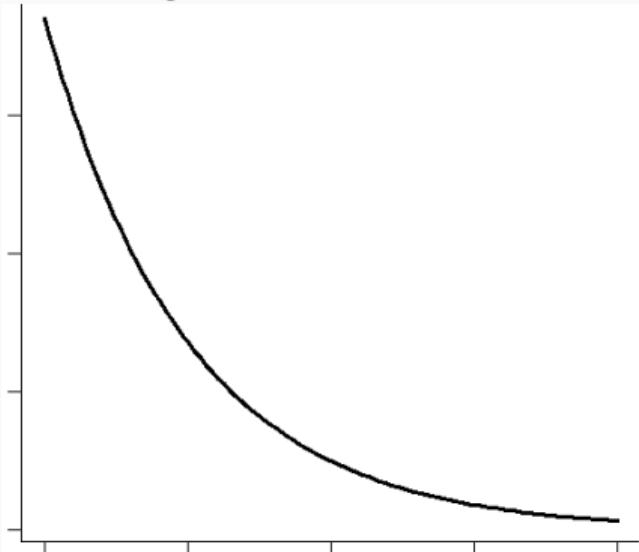
Our model assumes that an impression on day t has an effect on the number of transactions on day t and that those impressions have no effect on sales on other days.

Do we really believe that ads only work on the day they were served/sent?

Our intuition as marketers tells us this probably isn't true and our email holdout test suggests that emails last about 3 days.

Exponential decay of advertising

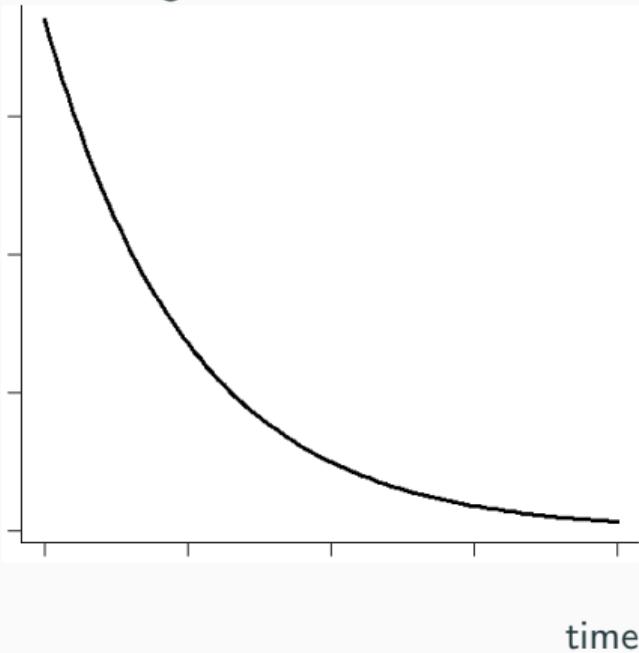
advertising effect



One theory of advertising response is that an ad had its biggest effect just after it is shown to the user and then the effect wears off over time.

Exponential decay of advertising

advertising effect

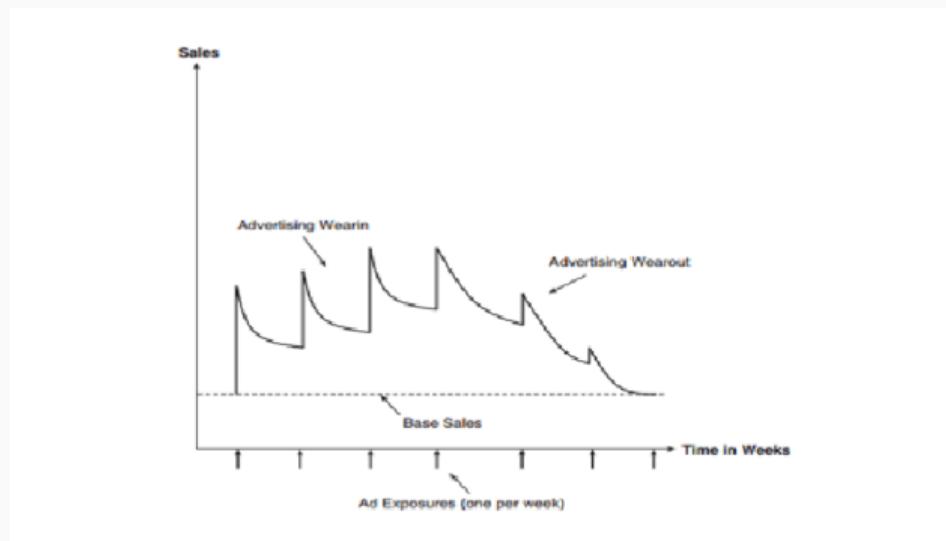


One theory of advertising response is that an ad had its biggest effect just after it is shown to the user and then the effect wears off over time.

We typically use an exponential decay function to describe how the effect of the ad falls off.

Ad stock variables in R

An ad stock variable is created by computing the exponential decay of the impressions on each day and then summing up the total “stock” from impressions on previous days. The result will look something like this:



Source: [Michael Wolfe](#)

Creating ad stock variables in R

Computing ad stock variables may sound complicated, but it is easy to do in R.

```
> mdata$email.stock <- as.numeric(filter(x=mdata$email, filter=0.5, method="recursive"))
> mdata$display.stock <- as.numeric(filter(x=mdata$display, filter=0.3, method="recursive"))
> mdata$direct.stock <- as.numeric(filter(x=mdata$direct, filter=0.75, method="recursive"))
> mdata$social.stock <- as.numeric(filter(x=mdata$social, filter=0.3, method="recursive"))
```

Ad stock variables

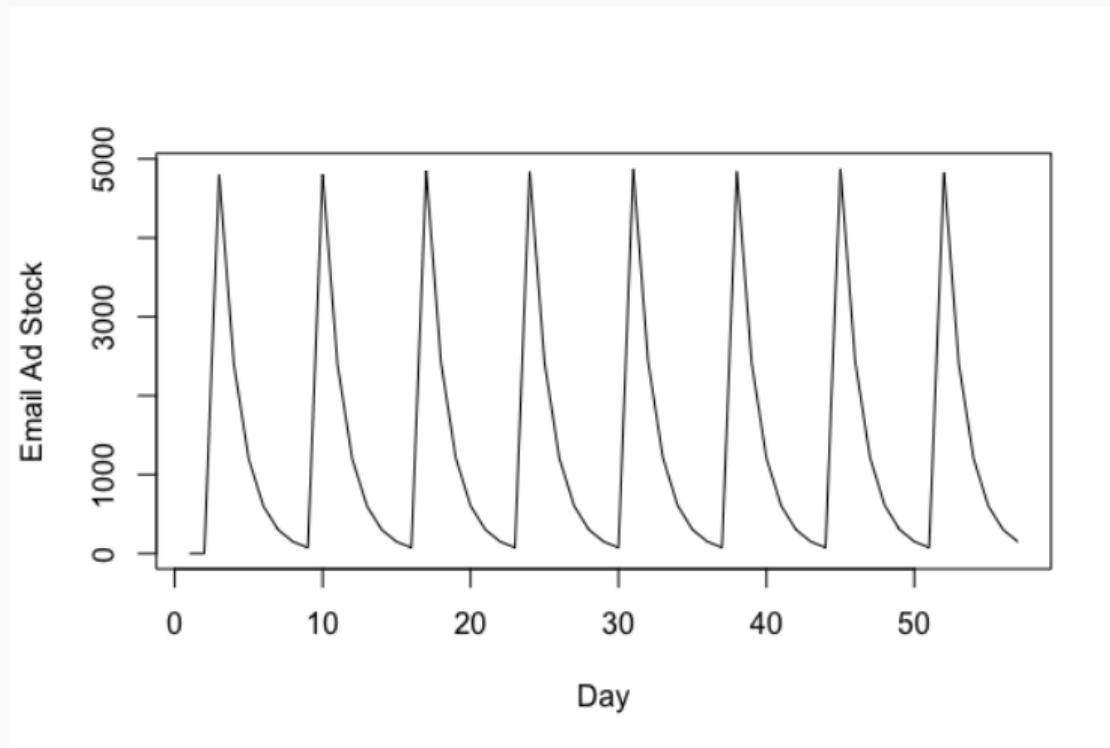
	trans	direct	display	email	email.holdout	social	dayofweek
2017-01-01	325	0	3786	0		7481	Sunday
2017-01-02	357	0	3792	0		7416	Monday
2017-01-03	589	0	3656	4798	1203	7505	Tuesday
2017-01-04	479	0	3731	0		7648	Wednesday
2017-01-05	403	0	3770	0		7620	Thursday
2017-01-06	498	4974	3611	0		7614	Friday
	email.stock	display.stock	direct.stock	social.stock			
2017-01-01	0.00	3786.000		0	7481.00		
2017-01-02	0.00	4927.800		0	9660.30		
2017-01-03	4798.00	5134.340		0	10403.09		
2017-01-04	2399.00	5271.302		0	10768.93		
2017-01-05	1199.50	5351.391		0	10850.68		
2017-01-06	599.75	5216.417	4974		10869.20		

This might be easier with a picture.

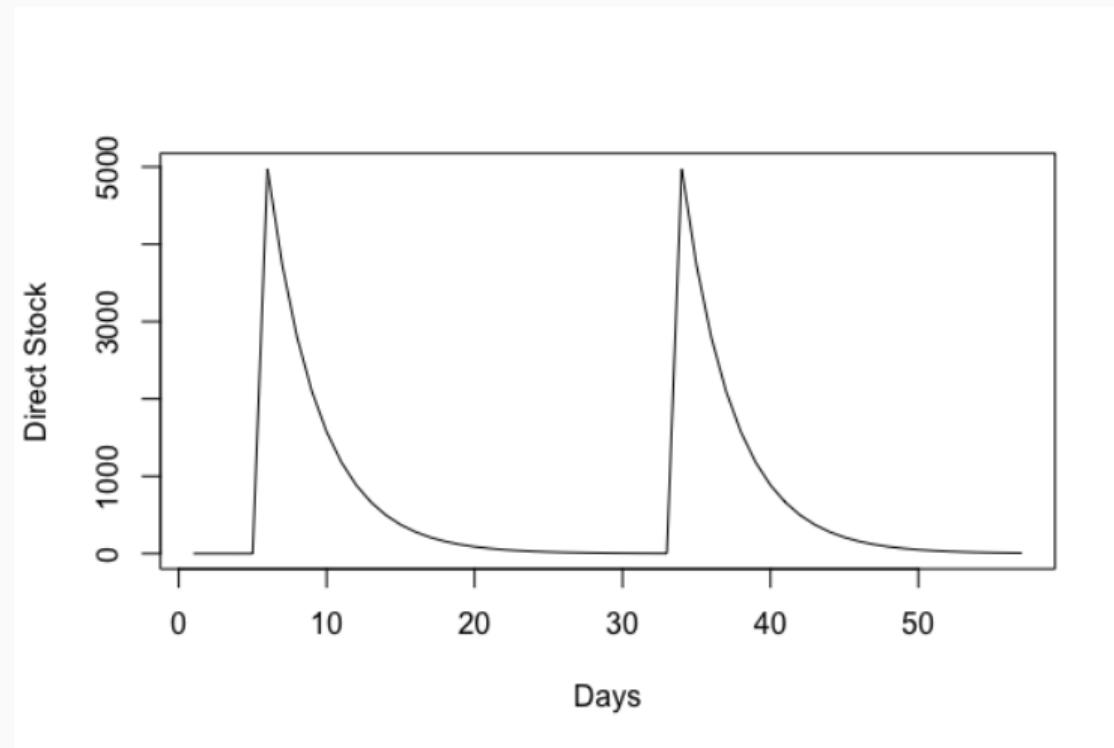
Create graphs of adstocks

```
> plot(mdata$email.stock, type="l", xlab="Day", ylab="Email Ad Stock")
> plot(mdata$display.stock, type="l", xlab="Day", ylab="Display Ad Stock")
> plot(mdata$direct.stock, type="l", xlab="Days", ylab="Direct Stock")
> plot(mdata$social.stock, type="l", xlab="Days", ylab="Social Stock")
```

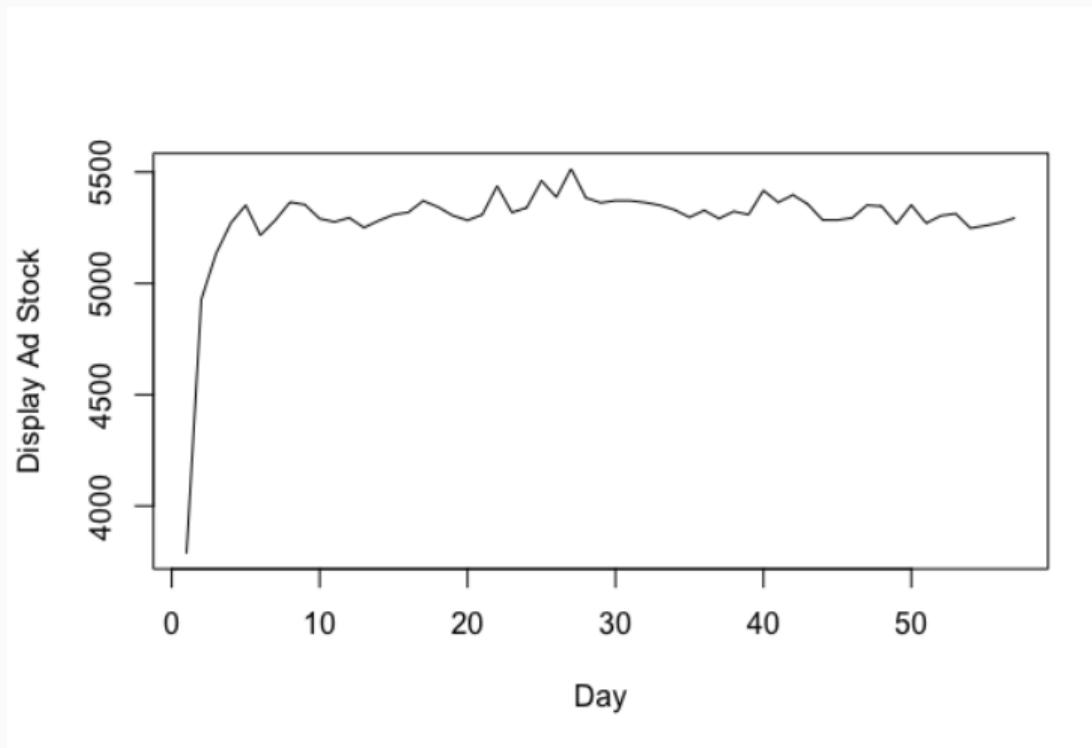
Email ad stock variable



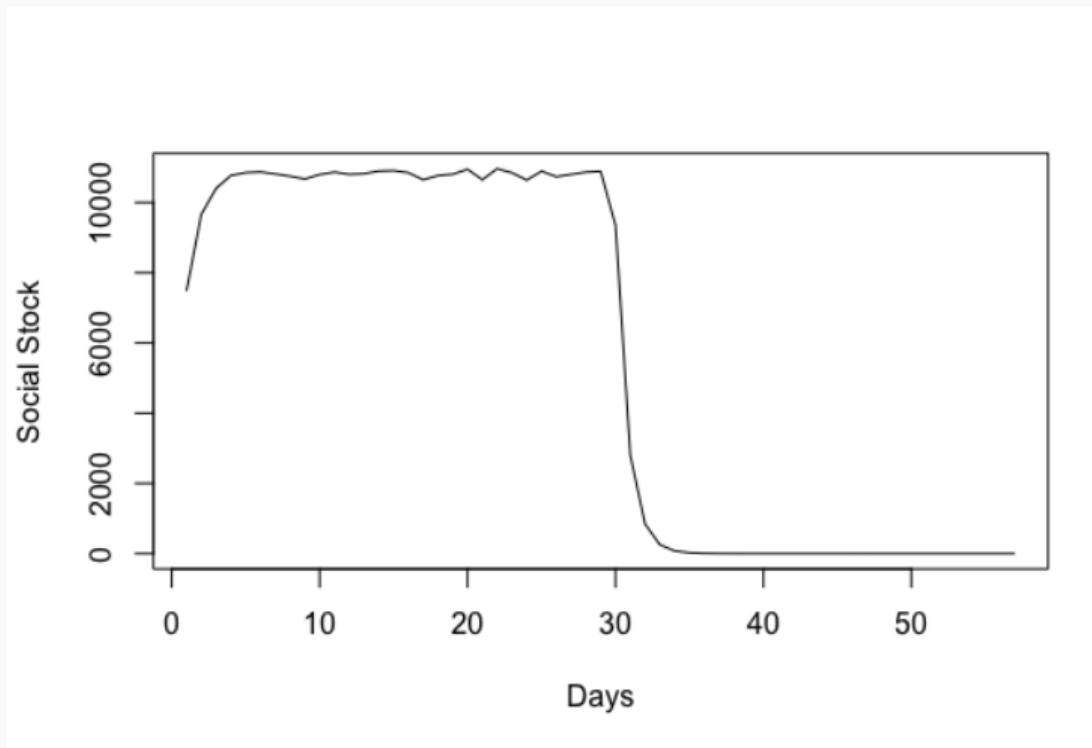
Direct ad stock variable



Display ad stock variable



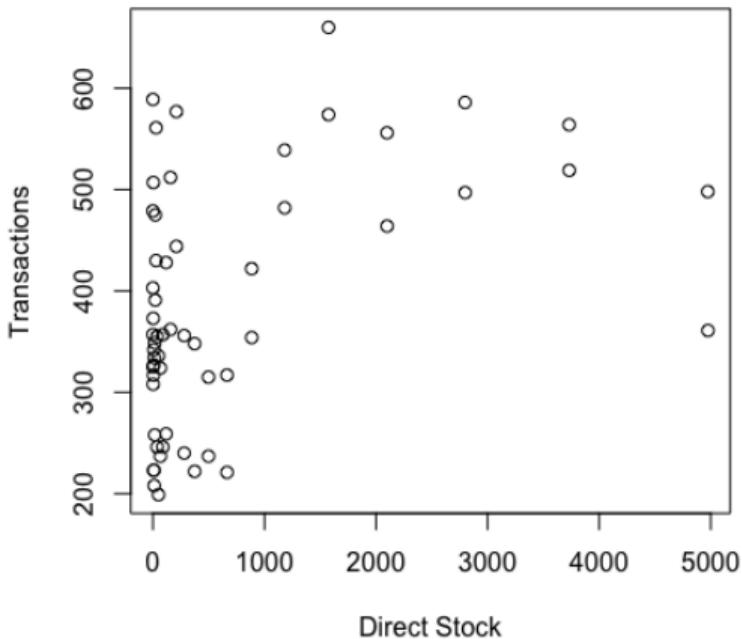
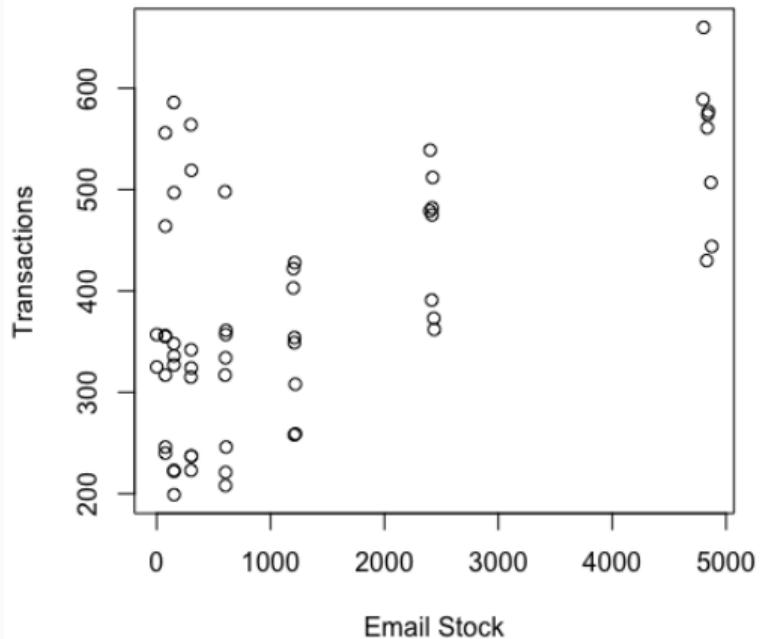
Social ad stock variable



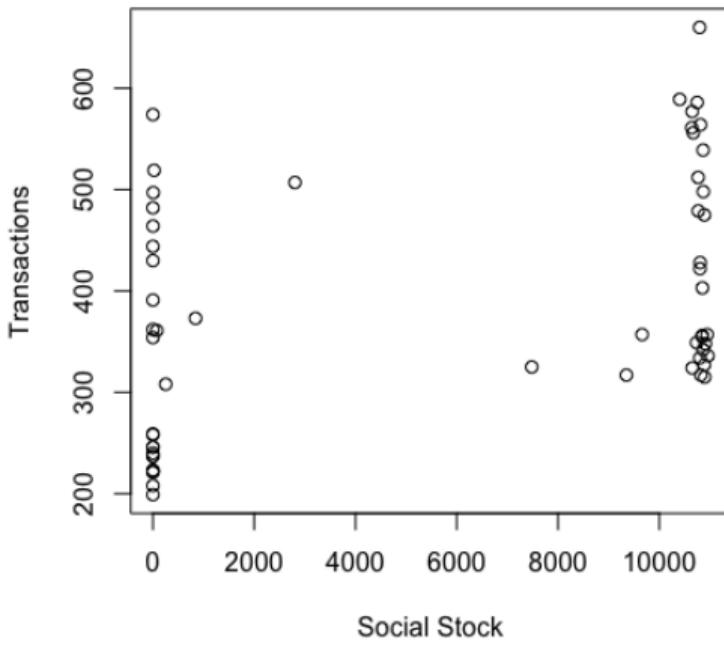
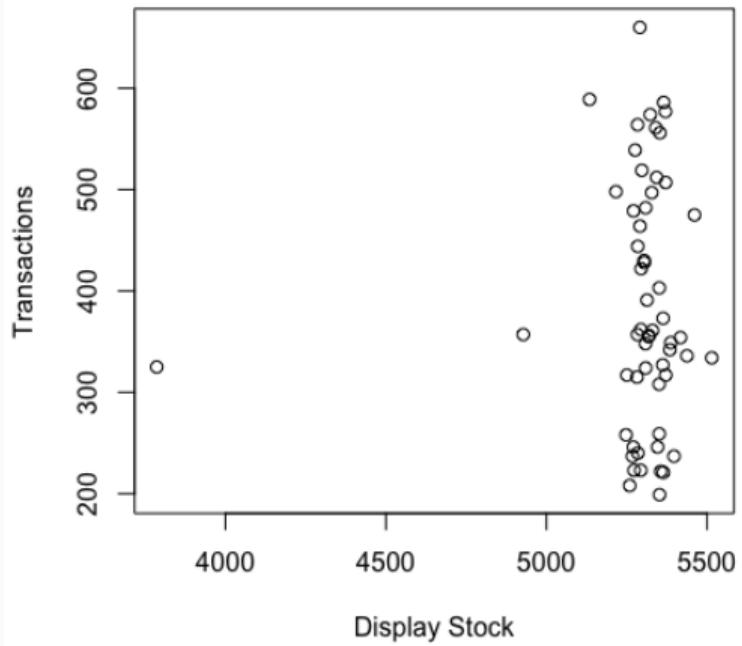
Plot ad stock variables against transactions

```
> plot(mdata$email.stock, type="l", xlab="Day", ylab="Email Ad Stock")
> plot(mdata$display.stock, type="l", xlab="Day", ylab="Display Ad Stock")
> plot(mdata$direct.stock, type="l", xlab="Days", ylab="Direct Stock")
> plot(mdata$social.stock, type="l", xlab="Days", ylab="Social Stock")
```

Email and direct ad stock versus transactions



Display and social ad stock versus transactions



Marketing mix model with ad stock variables

```
> m3 <- lm(trans~email.stock+display.stock+direct.stock+social.stock,
+           data=mdata[5:nrow(mdata),]) # Remove first few observations to allow for
> summary(m3)
...
Coefficients:
              Estimate Std. Error t value     Pr(>|t|)
(Intercept) -460.452735  753.753237  -0.611      0.544
email.stock    0.050538    0.004814   10.499 0.000000000000503 ***
display.stock   0.128309    0.141450    0.907      0.369
direct.stock    0.055573    0.006193    8.974 0.0000000000077486 ***
social.stock    0.009468    0.001422    6.659 0.0000000245161516 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 54.17 on 48 degrees of freedom
Multiple R-squared:  0.8147,    Adjusted R-squared:  0.7992
F-statistic: 52.74 on 4 and 48 DF,  p-value: < 0.0000000000000022
```

Things to notice about our third marketing mix model

```
> m3 <- lm(trans~email.stock+display.stock+direct.stock+social.stock,
+           data=mdata[5:nrow(mdata),])
> summary(m3)

Call:
lm(formula = trans ~ email.stock + display.stock + direct.stock +
    social.stock, data = mdata[5:nrow(mdata), ])

Residuals:
    Min      1Q  Median      3Q     Max 
-170.442 -31.630  -3.793  19.411 125.000 

Coefficients:
            Estimate Std. Error t value
Pr(>|t|)    
(Intercept) -460.452735  753.753237  -0.611
0.544
email.stock   0.050538   0.004814
10.499 0.0000000000000503 *** 
display.stock   0.128309   0.141450    0.907
0.369
direct.stock   0.055573   0.006193
8.974 0.00000000077486 *** 
social.stock    0.009468   0.001422
6.659 0.000000245161516 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 54.17 on 48 degrees of freedom
Multiple R-squared:  0.8147,    Adjusted R-squared:  0.7992 
F-statistic: 52.74 on 4 and 48 DF,
p-value: < 0.000000000000022
```

We now find positive effects for all forms of advertising.

Things to notice about our third marketing mix model

```
> m3 <- lm(trans~email.stock+display.stock+direct.stock+social.stock,
+           data=mdata[5:nrow(mdata),])
> summary(m3)

Call:
lm(formula = trans ~ email.stock + display.stock + direct.stock +
    social.stock, data = mdata[5:nrow(mdata), ])

Residuals:
    Min      1Q  Median      3Q     Max 
-170.442 -31.630   -3.793  19.411 125.000 

Coefficients:
            Estimate Std. Error t value
Pr(>|t|)    
(Intercept) -460.452735  753.753237 -0.611
0.544
email.stock    0.050538    0.004814
10.499 0.0000000000000503 ***
display.stock   0.128309    0.141450    0.907
0.369
direct.stock    0.055573    0.006193
8.974 0.00000000077486 ***
social.stock     0.009468    0.001422
6.659 0.000000245161516 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 54.17 on 48 degrees of freedom
Multiple R-squared:  0.8147,    Adjusted R-squared:  0.7992 
F-statistic: 52.74 on 4 and 48 DF,
p-value: < 0.0000000000000022
```

We now find positive effects for all forms of advertising.

Email and display seem to have similar effects on transactions at about 50 additional transactions on the first day the piece is sent.

Things to notice about our third marketing mix model

```
> m3 <- lm(trans~email.stock+display.stock+direct.stock+social.stock,
+           data=mdata[5:nrow(mdata),])
> summary(m3)

Call:
lm(formula = trans ~ email.stock + display.stock + direct.stock +
    social.stock, data = mdata[5:nrow(mdata), ])

Residuals:
    Min      1Q  Median      3Q     Max 
-170.442 -31.630  -3.793  19.411 125.000 

Coefficients:
            Estimate Std. Error t value
Pr(>|t|)    
(Intercept) -460.452735  753.753237 -0.611
0.544
email.stock   0.050538   0.004814
10.499 0.0000000000000503 ***
display.stock  0.128309   0.141450   0.907
0.369
direct.stock   0.055573   0.006193
8.974 0.00000000077486 ***
social.stock    0.009468   0.001422
6.659 0.000000245161516 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 54.17 on 48 degrees of freedom
Multiple R-squared:  0.8147,    Adjusted R-squared:  0.7992 
F-statistic: 52.74 on 4 and 48 DF,
p-value: < 0.0000000000000022
```

We now find positive effects for all forms of advertising.

Email and display seem to have similar effects on transactions at about 50 additional transactions on the first day the piece is sent.

All effects are statistically significant except for display. Display still has a high standard error because there is not much variation in the display stock. We simply can't measure the effect of display using this data.

Estimating the decay rate

In this example, we assumed a decay rate for each channel when we created the ad stock variables.

```
> mdata$email.stock <- as.numeric(filter(x=mdata$email, filter=0.5,  
                                         method="recursive"))  
> mdata$display.stock <- as.numeric(filter(x=mdata$display, filter=0.3,  
                                         method="recursive"))  
> mdata$direct.stock <- as.numeric(filter(x=mdata$direct, filter=0.75,  
                                         method="recursive"))  
> mdata$social.stock <- as.numeric(filter(x=mdata$social, filter=0.3,  
                                         method="recursive"))
```

I selected a short decay rate for display and social and longer ones for email and catalog based on intuition. More sophisticated approaches estimate the decay rate from the data.

Interactions between channels

An interaction occurs when there is an extra "boost" to having two advertising channels active at the same time. Together, the two channels are more effective than the sum of their parts. There is a synergy.

Interactions between channels

An interaction occurs when there is an extra "boost" to having two advertising channels active at the same time. Together, the two channels are more effective than the sum of their parts. There is a synergy.

We model this by adding an extra term in our model that is the multiple of two other variables.

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \beta_5 (\text{email} \times \text{direct}) + \epsilon_t$$

This is really easy to do in R.

Adding interactions to a model in R

```
> m4 <- lm(trans~email.stock+display.stock+direct.stock+social.stock+
+           email.stock*direct.stock, data=mdata[5:nrow(mdata),])
> summary(m4)
...
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -418.432368566 764.864071866 -0.547 0.587
email.stock   0.052249390 0.006017106  8.683 0.000000000248 ***
display.stock  0.120103084 0.143613018  0.836 0.407
direct.stock   0.057597635 0.007529130  7.650 0.0000000008521 ***
social.stock   0.009506492 0.001435581  6.622 0.0000000306335 ***
email.stock:direct.stock -0.000003204 0.000006659 -0.481 0.633
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 54.61 on 47 degrees of freedom
Multiple R-squared:  0.8156,    Adjusted R-squared:  0.7959
F-statistic: 41.57 on 5 and 47 DF,  p-value: 0.00000000000003867
```

There is no significant interaction effect between email and direct mail.

Is there more to the ad response story?

So far, we have explored a few advanced features to our model:

- **Ad Stock / Exponential decay**: Ads have an effect on transactions that decays over time
- **Interactions**: Ads may be synergistic

Is there more to the ad response story?

So far, we have explored a few advanced features to our model:

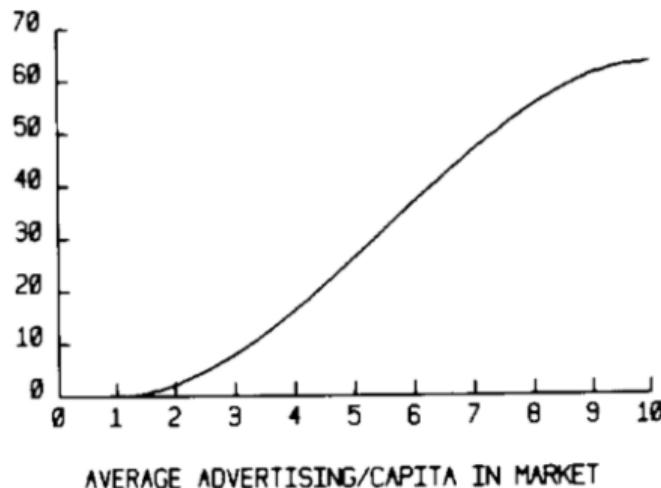
- **Ad Stock / Exponential decay**: Ads have an effect on transactions that decays over time
- **Interactions**: Ads may be synergistic

There are several other ideas about how advertising works that we could build into a marketing mix model, but are outside the scope of this workshop:

- **S-shaped advertising response**
- **Wearout**
- **Competitive advertising**

S Shaped advertising response

SALES/CAPITA



Two old saws of advertising:

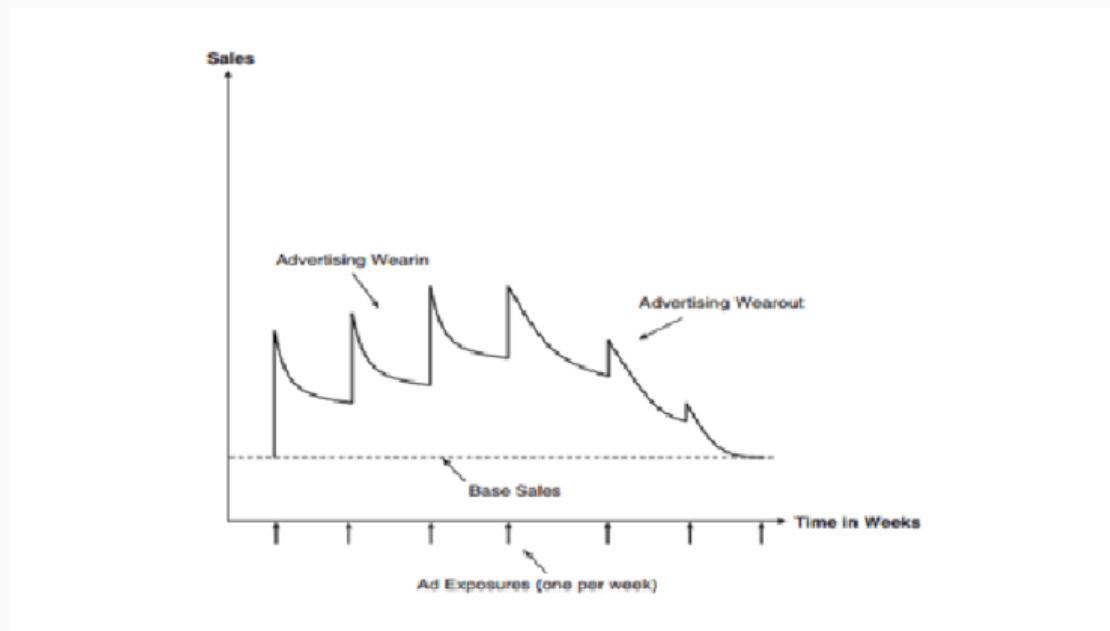
- Advertising doesn't have much effect on sales until you get enough "weight" in the market.
- There is a point of **diminishing returns** to advertising.

This can be modeled by assuming an "S-shaped" advertising response curve.

Source: Little (1979)

Wearin and wearout

Over time, even a constant level of advertising may become less effective, particularly if the creatives don't change.



Little's principles

Little (1979) summarized five key principals for marketing mix models, based on the data he had studied on aggregate advertising spending and sales.

The empirical evidence suggests that at least the following phenomena should be considered in building dynamic models of advertising response:

P1. Sales respond dynamically upward and downward to increases and decreases of advertising and frequently do so at different rates.

P2. Steady-state response can be concave or S-shaped and will often have positive sales at zero advertising.

P3. Competitive advertising affects sales.

P4. The dollar effectiveness of advertising can change over time as the result of changes in media, copy, and other factors.

P5. Products sometimes respond to increased advertising with a sales increase that falls off even as advertising is held constant.

Marketing mix modeling

What can go wrong?

Endogeneity bias

If the advertiser is buying more advertising ("heaving-up") during periods she knows have high demand (like holidays), there will be a correlation between advertising and sales.

Endogeneity bias

If the advertiser is buying more advertising (“heaving-up”) during periods she knows have high demand (like holidays), there will be a correlation between advertising and sales.

This correlation will show up in our regression and we might interpret it as “the effect of advertising”, when the causality is actually reversed.

Endogeneity bias

If the advertiser is buying more advertising (“heaving-up”) during periods she knows have high demand (like holidays), there will be a correlation between advertising and sales.

This correlation will show up in our regression and we might interpret it as “the effect of advertising”, when the causality is actually reversed.

Even though our final model (m3) seems to be giving us pretty good estimates of advertising response it might not be. This has the terrible name **endogeneity bias**.

Endogeneity bias corrections

There are some corrections for endogeneity bias in marketing mix models:

- Instrumental variables
- Model the process for setting advertising (e.g. VAR models)

Endogeneity bias corrections

There are some corrections for endogeneity bias in marketing mix models:

- Instrumental variables
- Model the process for setting advertising (e.g. VAR models)

These approaches have some major drawbacks and **holdout testing** is the only sure-fire way to accurately measure ad response.

Modeler degrees of freedom

As you have seen, there are many decisions for the modeler to make when building a marketing mix model.

- Does advertising last more than one day and if so, what is the **decay rate**?
- Should my model include **diminishing returns** to the level of advertising?
- Should my model include advertising **wearout**?
- Should I try to **correct endogeneity bias** and how?

Modeler degrees of freedom

As you have seen, there are many decisions for the modeler to make when building a marketing mix model.

- Does advertising last more than one day and if so, what is the **decay rate**?
- Should my model include **diminishing returns** to the level of advertising?
- Should my model include advertising **wearout**?
- Should I try to **correct endogeneity bias** and how?

The problem with having so many modeling choices is that it is tempting to keep trying different models until you find the one that you (or your client) like. This is sometimes called the **garden of forking paths** and it is a dangerous place for an analyst to be! You may end up with a model you like rather than one that is right.

Summary: marketing mix models

A **marketing mix model** is a **regression** relating advertising spending or total impressions to some response like sales.

To build a marketing mix model, you assemble data on advertising spending or impressions by day/week/month and conversions or sales by day/week/month and then fit a model using any regression tool.

The coefficients of this model tell us the correlations between daily/weekly/monthly levels of advertising and sales.

What can go wrong with marketing mix modeling?

There isn't enough variation in the advertising over time or two different types of advertising tend to happen at the same time.

- Model estimates will show large standard errors on the coefficients. You'll be stuck, but at least you'll know it.

What can go wrong with marketing mix modeling?

There isn't enough variation in the advertising over time or two different types of advertising tend to happen at the same time.

- Model estimates will show large standard errors on the coefficients. You'll be stuck, but at least you'll know it.

Advertising happens during periods of peak demand, reversing the causality.

- Model estimates may appear reasonable but will overestimate advertising response. You'll be wrong, but you won't know it.

What can go wrong with marketing mix modeling?

There isn't enough variation in the advertising over time or two different types of advertising tend to happen at the same time.

- Model estimates will show large standard errors on the coefficients. You'll be stuck, but at least you'll know it.

Advertising happens during periods of peak demand, reversing the causality.

- Model estimates may appear reasonable but will overestimate advertising response. You'll be wrong, but you won't know it.

The analyst can twiddle with the model until she gets it to say what she likes or what you want to hear.

Model-based attribution

Model-based attribution

What is model-based attribution?

Model-based attribution

Model-based attribution is similar to marketing mix modeling, but we do the analysis at the user-level, relating a user's transactions or conversions to her prior advertising exposures. A simple attribution model is:

$$\text{conversion}_{it} = \beta_0 + \beta_1 \text{display}_{it} + \beta_2 \text{social}_{it} + \beta_3 \text{email}_{it} + \beta_4 \text{direct}_{it} + \epsilon_{it}$$

The key difference from a marketing mix model is that now all the variables are indexed by i in addition to t .

Model-based attribution

Model-based attribution in R

Model-based attribution in R

To run this regression, we need to create a data frame where each row is a user-day and summarizes the user's impressions and transactions on that day. We have to transform the raw data to this:

```
> head(adata)
  id      date direct display email email.holdout social trans past.purchase has.email has.direct
1 1 2017-01-01     0      0          0      1      0          0      0          0      1
2 1 2017-01-02     0      0          0      1      0          0      0          0      1
3 1 2017-01-03     0      0          0      0      0          0      0          0      1
4 1 2017-01-04     0      0          0      0      0          0      0          0      1
5 1 2017-01-05     0      0          0      1      0          0      0          0      1
6 1 2017-01-06     1      0          0      0      2          0      0          0      1
```

My R code for doing this is a bit messy and takes several steps. I'm sure you could do better with [tidyverse](#).

Data transform step 1: Summary of impressions by user and day

First, we create a data frame that has a daily summary of each user's transactions.

```
> adatal <- as.data.frame(xtabs(~ id + date + channel, data=impress), stringsAsFactors=FALSE)
> adatal$id <- as.integer(adatal$id)
> adatal$date <- as.Date(adatal$date)
> adatal$channel <- as.factor(adatal$channel)
> dimnames(adatal)[[2]][4] <- "impr"
> head(adatal)
   id      date channel impr
1  1 2016-12-31  direct    0
2  2 2016-12-31  direct    0
3  3 2016-12-31  direct    0
4  5 2016-12-31  direct    0
5  6 2016-12-31  direct    0
6  7 2016-12-31  direct    0
```

Data transform step 2: Adding in users with zero impressions

Because some users don't appear in the impressions file, we have to add in rows (with zeros) for those users.

```
> pop <- unique(cust$id)
> no.impress.ids <- pop[!(pop %in% unique(impress$id))]
> dates <- sort(unique(impress$date))
> channels <- unique(impress$channel)
> no.impress.obs <- data.frame(id=rep(no.impress.ids, each=length(dates)*length(channels)),
+                                 date=rep(rep(dates, each=length(channels)), length(no.impress.ids)),
+                                 channel=rep(channels, length(no.impress.ids)*length(dates)),
+                                 impr=rep(0, length(dates)*length(no.impress.ids)*length(channels)),
+                                 stringsAsFactors=FALSE)
> no.impress.obs$channel <- as.factor(no.impress.obs$channel)
> adatal <- rbind(adatal, no.impress.obs)
> summary(adatal)
      id          date           channel        impr
Min. : 1  Min. :2016-12-31  direct     :590000  Min. : 0.0000
1st Qu.: 2501  1st Qu.:2017-01-14  display    :590000  1st Qu.: 0.0000
Median : 5000  Median :2017-01-29  email      :590000  Median : 0.0000
Mean   : 5000  Mean   :2017-01-29  email.holdout:590000  Mean   : 0.1699
3rd Qu.: 7500  3rd Qu.:2017-02-13  social     :590000  3rd Qu.: 0.0000
Max.   :10000  Max.   :2017-02-27                    Max.   :16.0000
```

Data transform step 3: Switch to wide format

Next, we unstack the impressions column by the channel.

```
> adata <- reshape(adatal, direction="wide", v.names="impr", idvar=c("id", "date"),
+                     timevar="channel", new.row.names=NULL)
> sum(adata$impr.direct) == length(impress$channel[impress$channel=="direct"]) #quick check
[1] TRUE
> nrow(adata)
[1] 590000
> head(adata)
   id      date impr.direct impr.display impr.email impr.email.holdout impr.social
1  1 2016-12-31          0          0          0              0              0
2  2 2016-12-31          0          0          0              0              0
3  3 2016-12-31          0          0          0              0              0
4  5 2016-12-31          0          1          0              0              0
5  6 2016-12-31          0          0          0              0              0
6  7 2016-12-31          0          0          0              0              0
```

Data transformation step 4: Add the daily transactions for each user

Finally, we summarize the transactions by customer-day and merge with impressions.

```
> atrans <- as.data.frame(xtabs(~ id + date, data=trans), stringsAsFactors=FALSE)
> atrans$id <- as.integer(atrans$id)
> atrans$date <- as.Date(atrans$date)
> dimnames(atrans)[[2]][3] <- "trans"
> adata <- merge(adata, atrans, by=c("id", "date"), all=TRUE)
> adata$trans[is.na(adata$trans)] <- 0 # fill in zeros for transactions
> head(adata)
   id      date impr.direct impr.display impr.email impr.email.holdout impr.social trans
1  1 2016-12-31          0          0          0            0            0            0
2  1 2017-01-01          0          0          0            0            1            0
3  1 2017-01-02          0          0          0            0            1            0
4  1 2017-01-03          0          0          0            0            0            0
5  1 2017-01-04          0          0          0            0            0            0
6  1 2017-01-05          0          0          0            0            1            0
```

Data transformation step 5: Final tidy up of attribution modeling data

```
> # Remove first and last days (which are incomplete)
> adata <- adata[adata$date!="2016-12-31" & adata$date != "2017-02-28" & adata$date != "2017-02-27",]
# Add customer info from cust table
> adata <- merge(adata, cust, by=c("id"))
# Tidy up column names
> dimnames(adata)[[2]][3:11] <- c("direct", "display", "email", "email.holdout", "social",
   "trans", "past.purchase", "has.email", "has.direct")
> rm(adatal, atrans)
```

Summary of attribution modeling data

```
> summary(adata)
      id          date        direct       display       email
Min. : 1   Min. :2017-01-01   Min. :0.00000   Min. : 0.0000   Min. :0.00000
1st Qu.: 2501 1st Qu.:2017-01-15  1st Qu.:0.00000  1st Qu.: 0.0000  1st Qu.:0.00000
Median : 5000 Median :2017-01-29  Median :0.00000  Median : 0.0000  Median :0.00000
Mean   : 5000 Mean  :2017-01-29  Mean   :0.01745  Mean   : 0.3731  Mean   :0.06741
3rd Qu.: 7500 3rd Qu.:2017-02-12 3rd Qu.:0.00000 3rd Qu.: 0.0000  3rd Qu.:0.00000
Max.  :10000 Max.  :2017-02-26  Max.  :1.00000  Max.  :13.0000  Max.  :1.00000
email.holdout social        trans      past.purchase has.email
Min. :0.00000  Min. : 0.0000  Min. :0.00000  Min. : 0.0000  Min. :0.0000
1st Qu.:0.00000 1st Qu.: 0.0000  1st Qu.:0.00000  1st Qu.: 0.0000  1st Qu.:0.0000
Median :0.00000 Median : 0.0000  Median :0.00000  Median :1.0000  Median :1.0000
Mean   :0.01681 Mean  : 0.3954  Mean   :0.03858  Mean   : 0.5022  Mean   :0.6001
3rd Qu.:0.00000 3rd Qu.: 0.0000  3rd Qu.:0.00000  3rd Qu.:1.0000  3rd Qu.:1.0000
Max.  :1.00000 Max.  :16.0000  Max.  :1.00000  Max.  :1.0000  Max.  :1.0000
has.direct
Min. :0.0000
1st Qu.:0.0000
Median :0.0000
Mean   :0.4974
3rd Qu.:1.0000
Max.  :1.0000
```

Summary of attribution modeling data on a specific day

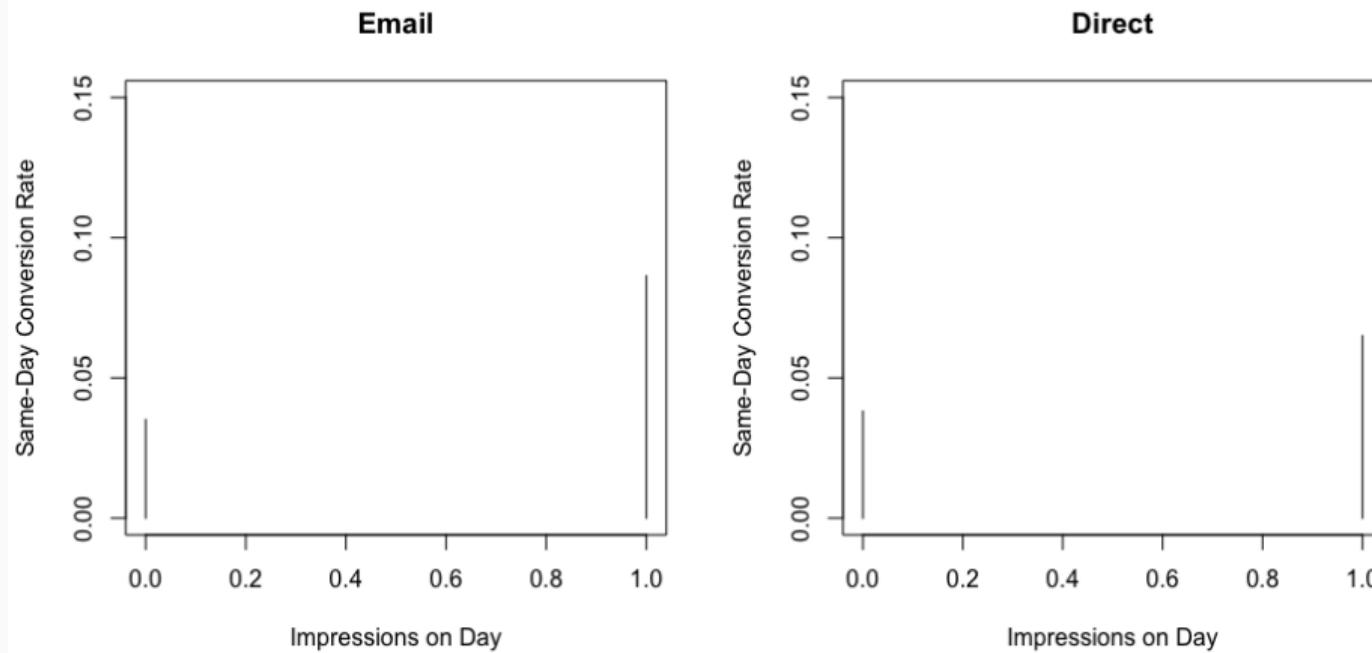
```
> summary(adata[adata$date=="2017-01-03",])
    id          date        direct      display       email     email.holdout
Min. : 1   Min. :2017-01-03  Min. :0   Min. :0.0000  Min. :0.0000  Min. :0.0000
1st Qu.: 2501  1st Qu.:2017-01-03  1st Qu.:0   1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.:0.0000
Median : 5000  Median :2017-01-03  Median :0   Median :0.0000  Median :0.0000  Median :0.0000
Mean   : 5000  Mean   :2017-01-03  Mean   :0   Mean   :0.3656  Mean   :0.4798  Mean   :0.1203
3rd Qu.: 7500  3rd Qu.:2017-01-03  3rd Qu.:0   3rd Qu.:0.0000  3rd Qu.:1.0000  3rd Qu.:0.0000
Max.   :10000  Max.   :2017-01-03  Max.   :0   Max.   :9.0000  Max.   :1.0000  Max.   :1.0000

    social        trans      past.purchase has.email      has.direct
Min. : 0.0000  Min. :0.0000  Min. :0.0000  Min. :0.0000  Min. :0.0000
1st Qu.: 0.0000  1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.:0.0000
Median : 0.0000  Median :0.0000  Median :1.0000  Median :1.0000  Median :0.0000
Mean   : 0.7505  Mean   :0.0589  Mean   :0.5022  Mean   :0.6001  Mean   :0.4974
3rd Qu.: 1.0000  3rd Qu.:0.0000  3rd Qu.:1.0000  3rd Qu.:1.0000  3rd Qu.:1.0000
Max.   :12.0000  Max.   :1.0000  Max.   :1.0000  Max.   :1.0000  Max.   :1.0000
```

Visualizing the relationship between impressions and transactions

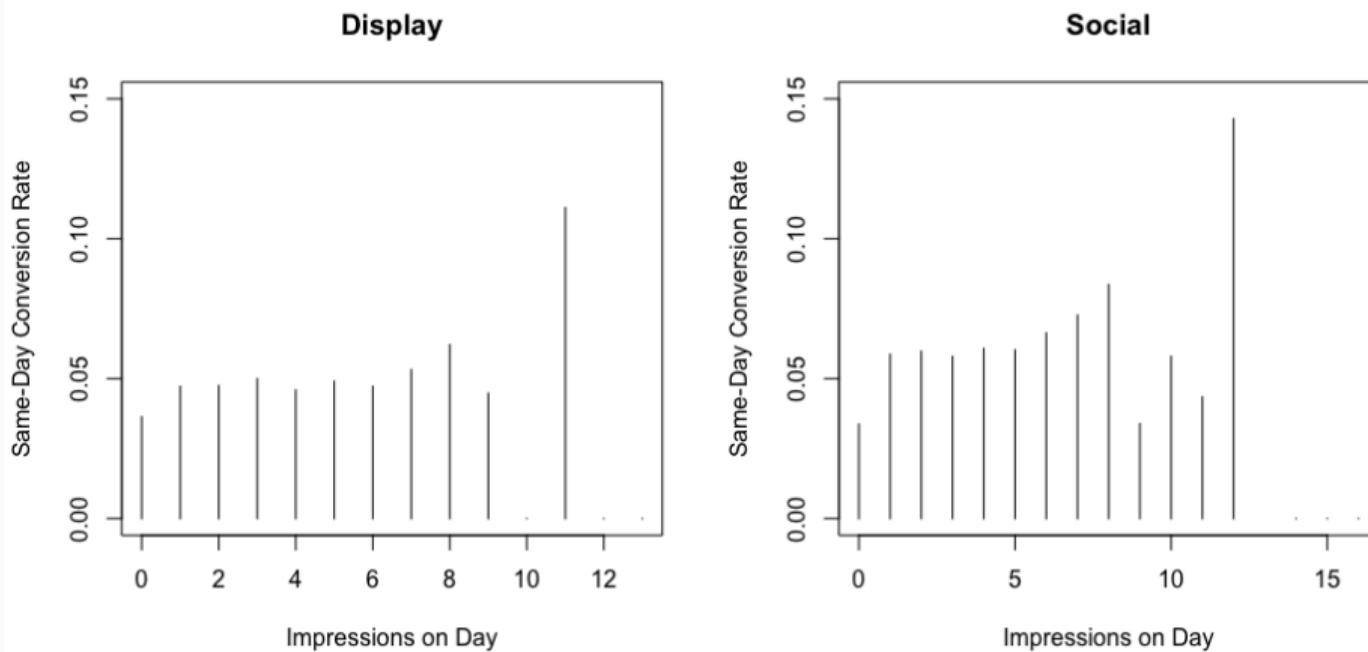
```
> plot(aggregate(trans~direct, data=adata, FUN=mean), type="h", ylim=c(0,0.15),
+       xlab="Impressions on Day", ylab="Same-Day Conversion Rate", main="Direct")
> plot(aggregate(trans~email, data=adata, FUN=mean), type="h", ylim=c(0,0.15),
+       xlab="Impressions on Day", ylab="Same-Day Conversion Rate", main="Email")
> plot(aggregate(trans~display, data=adata, FUN=mean), type="h", ylim=c(0,0.15),
+       xlab="Impressions on Day", ylab="Same-Day Conversion Rate", main="Display")
> plot(aggregate(trans~social, data=adata, FUN=mean), type="h", ylim=c(0,0.15),
+       xlab="Impressions on Day", ylab="Same-Day Conversion Rate", main="Social")
```

Direct and email daily impressions versus transactions



Looking at user-level data, it is easy to see that users convert more on days they get emails or direct mail.

Display and social daily impressions versus transactions



A simple attribution model

```
> m1 <- lm(trans ~ direct + display + email + social, data=adata)
> summary(m1)
...
Coefficients:
            Estimate Std. Error t value      Pr(>|t|)
(Intercept) 0.0296060  0.0003024   97.91 <0.0000000000000002 ***
direct       0.0306548  0.0019423   15.78 <0.0000000000000002 ***
display      0.0041618  0.0002759   15.09 <0.0000000000000002 ***
email        0.0520472  0.0010144   51.31 <0.0000000000000002 ***
social       0.0085516  0.0002542   33.64 <0.0000000000000002 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1919 on 569995 degrees of freedom
Multiple R-squared:  0.007249,  Adjusted R-squared:  0.007242
F-statistic: 1041 on 4 and 569995 DF,  p-value: < 0.0000000000000002
```

Key things to notice about our attribution model

```
> m1 <- lm(trans ~ direct + display + email + social, data=adata)
> summary(m1)
...
Coefficients:
            Estimate Std. Error t value     Pr(>|t|)
(Intercept) 0.0296060  0.0003024   97.91 <0.0000000000000002 ***
direct       0.0306548  0.0019423   15.78 <0.0000000000000002 ***
display      0.0041618  0.0002759   15.09 <0.0000000000000002 ***
email        0.0520472  0.0010144   51.31 <0.0000000000000002 ***
social       0.0085516  0.0002542   33.64 <0.0000000000000002 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1919 on 569995 degrees of freedom
Multiple R-squared:  0.007249,  Adjusted R-squared:  0.007242
F-statistic: 1041 on 4 and 569995 DF,
p-value: < 0.000000000000022
```

We have gotten reasonable estimates of advertising response even for this simple model that only looks at same-day response. For example, 1 email impressions is associated with 0.052 additional transactions on the same day.

Key things to notice about our attribution model

```
> m1 <- lm(trans ~ direct + display + email + social, data=adata)
> summary(m1)
...
Coefficients:
            Estimate Std. Error t value     Pr(>|t|)
(Intercept) 0.0296060  0.0003024   97.91 <0.0000000000000002 ***
direct       0.0306548  0.0019423   15.78 <0.0000000000000002 ***
display      0.0041618  0.0002759   15.09 <0.0000000000000002 ***
email        0.0520472  0.0010144   51.31 <0.0000000000000002 ***
social       0.0085516  0.0002542   33.64 <0.0000000000000002 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1919 on 569995 degrees of freedom
Multiple R-squared:  0.007249, Adjusted R-squared:  0.007242
F-statistic: 1041 on 4 and 569995 DF,
p-value: < 0.000000000000022
```

We have gotten reasonable estimates of advertising response even for this simple model that only looks at same-day response. For example, 1 email impressions is associated with 0.052 additional transactions on the same day.

The probability of transacting on a day where the user is exposed to no advertising is 3%. It would be bad to assume that all sales are due to advertising.

A model that includes user-characteristics

Because we have a user-level model, we can bring in user characteristics.

```
> m2 <- lm(trans ~ direct + display + email + social + past.purchase, data=adata)
> summary(m2)

Coefficients:
              Estimate Std. Error t value     Pr(>|t|)    
(Intercept) 0.0142352  0.0003889   36.60 <0.000000000000002 *** 
direct       0.0203596  0.0019427   10.48 <0.000000000000002 *** 
display      0.0041846  0.0002749   15.22 <0.000000000000002 *** 
email        0.0433566  0.0010205   42.49 <0.000000000000002 *** 
social        0.0086049  0.0002533   33.97 <0.000000000000002 *** 
past.purchase 0.0320725  0.0005130   62.52 <0.000000000000002 *** 
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1912 on 569994 degrees of freedom
Multiple R-squared:  0.01401,    Adjusted R-squared:  0.014 
F-statistic:  1620 on 5 and 569994 DF,  p-value: < 0.0000000000000022
```

A model that includes user-characteristics

```
> m2 <- lm(trans ~ direct + display + email + social +
>           past.purchase, data=adata)
> summary(m2)

Coefficients:
            Estimate Std. Error t value     Pr(>|t|)
(Intercept) 0.0142352  0.0003889   36.60 <0.0000000000000002 ***
direct       0.0203596  0.0019427   10.48 <0.0000000000000002 ***
display      0.0041846  0.0002749   15.22 <0.0000000000000002 ***
email        0.0433566  0.0010205   42.49 <0.0000000000000002 ***
social       0.0086049  0.0002533   33.97 <0.0000000000000002 ***
past.purchase 0.0320725  0.0005130   62.52 <0.0000000000000002 ***

---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1912 on 569994 degrees of freedom
Multiple R-squared:  0.01401,    Adjusted R-squared:  0.014
F-statistic: 1620 on 5 and 569994 DF,
p-value: < 0.0000000000000022
```

The model suggests that customers who have not made a purchase before are much less likely to make a purchase (1.4% versus 4.6% baseline rate without advertising).

A model that includes user-characteristics

```
> m2 <- lm(trans ~ direct + display + email + social +
>           past.purchase, data=adata)
> summary(m2)
Coefficients:
            Estimate Std. Error t value     Pr(>|t|)
(Intercept) 0.0142352  0.0003889   36.60 <0.0000000000000002 ***
direct       0.0203596  0.0019427   10.48 <0.0000000000000002 ***
display      0.0041846  0.0002749   15.22 <0.0000000000000002 ***
email        0.0433566  0.0010205   42.49 <0.0000000000000002 ***
social       0.0086049  0.0002533   33.97 <0.0000000000000002 ***
past.purchase 0.0320725  0.0005130   62.52 <0.0000000000000002 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1912 on 569994 degrees of freedom
Multiple R-squared:  0.01401,    Adjusted R-squared:  0.014
F-statistic: 1620 on 5 and 569994 DF,
p-value: < 0.0000000000000022
```

The model suggests that customers who have not made a purchase before are much less likely to make a purchase (1.4% versus 4.6% baseline rate without advertising).

The estimates of advertising effects are smaller than in our first model. This is because people who have made a purchase are more likely to be exposed to advertising. When we control for past purchase, we can see that the effects of advertising are smaller than we might have thought based on the first model.

A model that includes user-characteristics

```
> m2 <- lm(trans ~ direct + display + email + social +
>           past.purchase, data=adata)
> summary(m2)
Coefficients:
            Estimate Std. Error t value     Pr(>|t|)
(Intercept) 0.0142352  0.0003889   36.60 <0.0000000000000002 ***
direct       0.0203596  0.0019427   10.48 <0.0000000000000002 ***
display      0.0041846  0.0002749   15.22 <0.0000000000000002 ***
email        0.0433566  0.0010205   42.49 <0.0000000000000002 ***
social       0.0086049  0.0002533   33.97 <0.0000000000000002 ***
past.purchase 0.0320725  0.0005130   62.52 <0.0000000000000002 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1912 on 569994 degrees of freedom
Multiple R-squared:  0.01401,    Adjusted R-squared:  0.014
F-statistic: 1620 on 5 and 569994 DF,
p-value: < 0.0000000000000022
```

The model suggests that customers who have not made a purchase before are much less likely to make a purchase (1.4% versus 4.6% baseline rate without advertising).

The estimates of advertising effects are smaller than in our first model. This is because people who have made a purchase are more likely to be exposed to advertising. When we control for past purchase, we can see that the effects of advertising are smaller than we might have thought based on the first model.

The controls you choose to add to the model can be critical.

Logistic regression

Those of you who learned regression from a careful teacher may have noticed that I was violating a key assumption of the linear regression model.

Logistic regression

Those of you who learned regression from a careful teacher may have noticed that I was violating a key assumption of the linear regression model.

Our outcome variable is binary and so it would be better to use a logistic regression for binary outcomes.

Logistic regression for attribution

```
> m3 <- glm(trans ~ direct + display + email + social + past.purchase ,  
+           family=binomial() , data=adata)  
> summary(m3)  
...  
Coefficients:  


|               | Estimate  | Std. Error | z value  | Pr(> z )                |
|---------------|-----------|------------|----------|-------------------------|
| (Intercept)   | -4.022893 | 0.014013   | -287.078 | <0.0000000000000002 *** |
| direct        | 0.412554  | 0.041688   | 9.896    | <0.0000000000000002 *** |
| display       | 0.100630  | 0.006581   | 15.291   | <0.0000000000000002 *** |
| email         | 0.762355  | 0.020010   | 38.099   | <0.0000000000000002 *** |
| social        | 0.180079  | 0.005356   | 33.621   | <0.0000000000000002 *** |
| past.purchase | 0.957207  | 0.015672   | 61.079   | <0.0000000000000002 *** |

  
---  
Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1  
  
(Dispersion parameter for binomial family taken to be 1)  
  
Null deviance: 186297 on 569999 degrees of freedom  
Residual deviance: 178869 on 569994 degrees of freedom  
AIC: 178881
```

Logistic regression for attribution

```
> m3 <- glm(trans ~ direct + display + email + social + past.purchase,
+             family=binomial(), data=adata)
> summary(m3)

...
Coefficients:
            Estimate Std. Error z value      Pr(>|z|)
(Intercept) -4.022893   0.014013 -287.078 <0.0000000000000002 ***
direct       0.412554   0.041688    9.896 <0.0000000000000002 ***
display      0.100630   0.006581   15.291 <0.0000000000000002 ***
email        0.762355   0.020010   38.099 <0.0000000000000002 ***
social       0.180079   0.005356   33.621 <0.0000000000000002 ***
past.purchase 0.957207   0.015672   61.079 <0.0000000000000002 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 186297  on 569999  degrees of freedom
Residual deviance: 178869  on 569994  degrees of freedom
AIC: 178881

Number of Fisher Scoring iterations: 6
```

While this model is better, the logistic regression coefficients can be difficult to interpret.

Logistic regression for attribution

```
> m3 <- glm(trans ~ direct + display + email + social + past.purchase,
+             family=binomial(), data=adata)
> summary(m3)

Coefficients:
            Estimate Std. Error z value      Pr(>|z|)
(Intercept) -4.022893   0.014013 -287.078 <0.0000000000000002 ***
direct       0.412554   0.041688    9.896 <0.0000000000000002 ***
display      0.100630   0.006581   15.291 <0.0000000000000002 ***
email        0.762355   0.020010   38.099 <0.0000000000000002 ***
social       0.180079   0.005356   33.621 <0.0000000000000002 ***
past.purchase 0.957207   0.015672   61.079 <0.0000000000000002 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 186297  on 569999  degrees of freedom
Residual deviance: 178869  on 569994  degrees of freedom
AIC: 178881

Number of Fisher Scoring iterations: 6
```

While this model is better, the logistic regression coefficients can be difficult to interpret.

We are still seeing that direct and email have a bigger effect than a single display or social impression.

Model-based attribution

Advanced attribution model features

Advanced attribution modeling features

- Ad stock decay
 - Similar to marketing mix models
- Wearout
 - Can model reduced effectiveness of ads at the user level
- Interactions between channels
 - Similar to marketing mix models
- User differences in advertising response (**heterogeneity**)
 - Individual-level parameters for ad response can be used to target customers
- Endogeneity bias corrections
 - More and better options than in marketing mix

Advanced attribution modeling features

- Ad stock decay
 - Similar to marketing mix models
- Wearout
 - Can model reduced effectiveness of ads at the user level
- Interactions between channels
 - Similar to marketing mix models
- User differences in advertising response (heterogeneity)
 - Individual-level parameters for ad response can be used to target customers
- Endogeneity bias corrections
 - More and better options than in marketing mix

These are all things you should ask about when assessing an attribution provider. Many commercial attribution solutions are lacking one or more of these features.

User differences in advertising response (heterogeneity)

The models we have estimated have assumed that the effect of advertising is the same for all users. This can't possibly be true.

User differences in advertising response (heterogeneity)

The models we have estimated have assumed that the effect of advertising is the same for all users. This can't possibly be true.

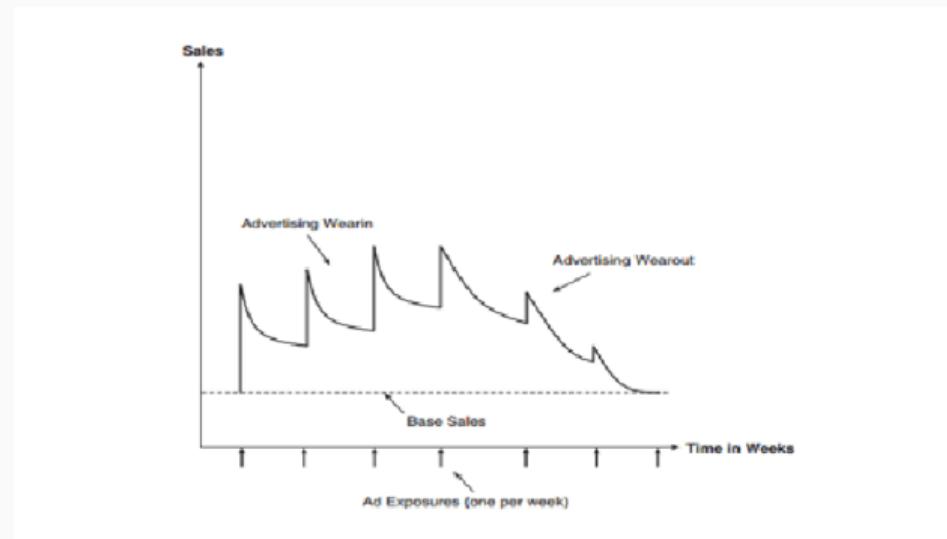
More advanced attribution models allow for differences between users in advertising response.

$$\text{conversion}_{it} = \beta_{0t} + \beta_{1t}\text{display}_{it} + \beta_{2t}\text{social}_{it} + \beta_{3t}\text{email} + \beta_{4t}\text{direct} + \epsilon_{it}$$

The key addition here is the t subscript on the β s.

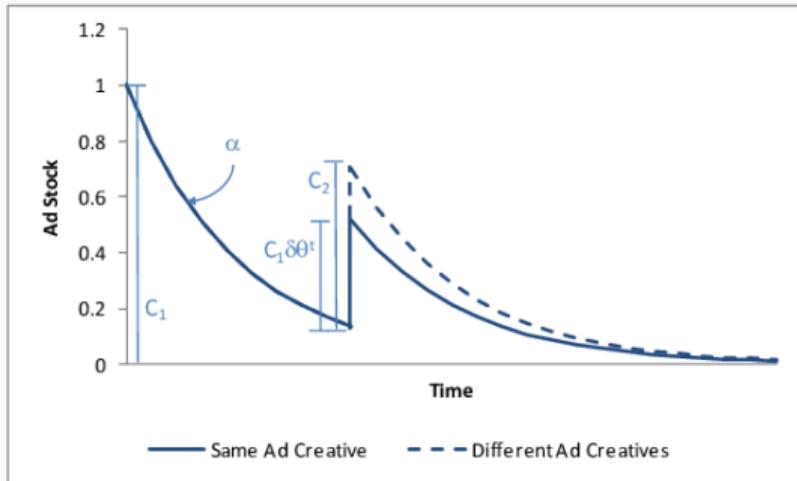
Ad stock at the individual-level

The ad stock concept makes a lot of sense when it is computed at the user-level. Goodwill toward the brand is a function of the previous impression that *an individual user* has received.



Source: [Michael Wolfe](#)

Effects of individual creatives: wearout



Source: Working version of [Braun and Moe \(2013\)](#)

Focusing on display ads only, Moe and Braun (2013) propose a model that includes:

- Different effects for each creative
- Wearout effects: the effect of a creative is a function of how long it has been since the user has last seen that creative

This type of model can be used to plan ad rotations.

This model is described in detail in [Braun and Moe \(2013\) Online Display Advertising: Modeling the Effects of Multiple Creatives and Individual Impression Histories](#).

Endogeneity bias corrections in attribution models

We have more and better options for endogeneity bias corrections in attribution models:

- Bring good user-level control variables into the model.
- Select a set of non-exposed users who “look similar” to those who were exposed (e.g. **propensity score matching**).
- If there was a sharp cutoff for those who did and did not receive the advertising, you can compare those just above to those just below the threshold (called **regression discontinuity**).

You should ask attribution vendors to explain their approach to correcting endogeneity bias.

Model-based attribution

**Example: Zantedeschi, Feit and
Bradlow (2017)**

Example: catalog and email response for an online retailer



- 300 customers from the the retailer's CRM system
 - E-mails sent to each customer
 - Catalogs sent to each customer
 - Daily purchase amount (direct/in-store) for each customer
 - > 90% of transactions can be matched to an existing customer by name, credit card, etc.
- Observations from 1 April 2012 to 10 May 2014
- All of these customers were targeted in each campaign, but sometimes were included in a holdout test

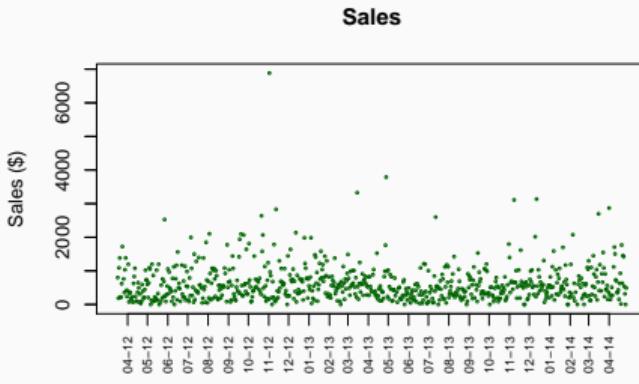
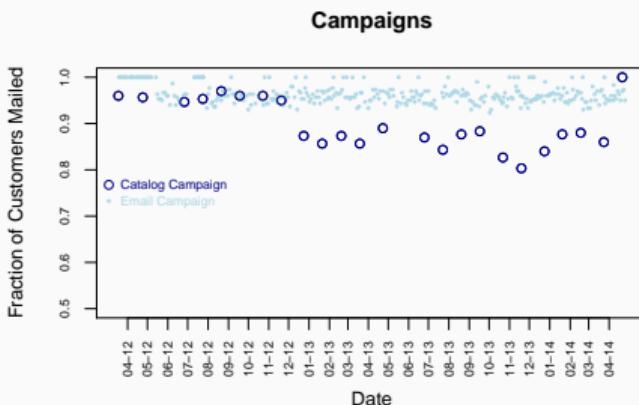
Example: catalog and email response for an online retailer



- 300 customers from the the retailer's CRM system
 - E-mails sent to each customer
 - Catalogs sent to each customer
 - Daily purchase amount (direct/in-store) for each customer
 - > 90% of transactions can be matched to an existing customer by name, credit card, etc.
- Observations from 1 April 2012 to 10 May 2014
- All of these customers were targeted in each campaign, but sometimes were included in a holdout test

This attribution model is reported in [Zantedeschi, Feit and Bradlow \(2017\) Measuring multi-channel advertising response, *Management Science*, forthcoming.](#)

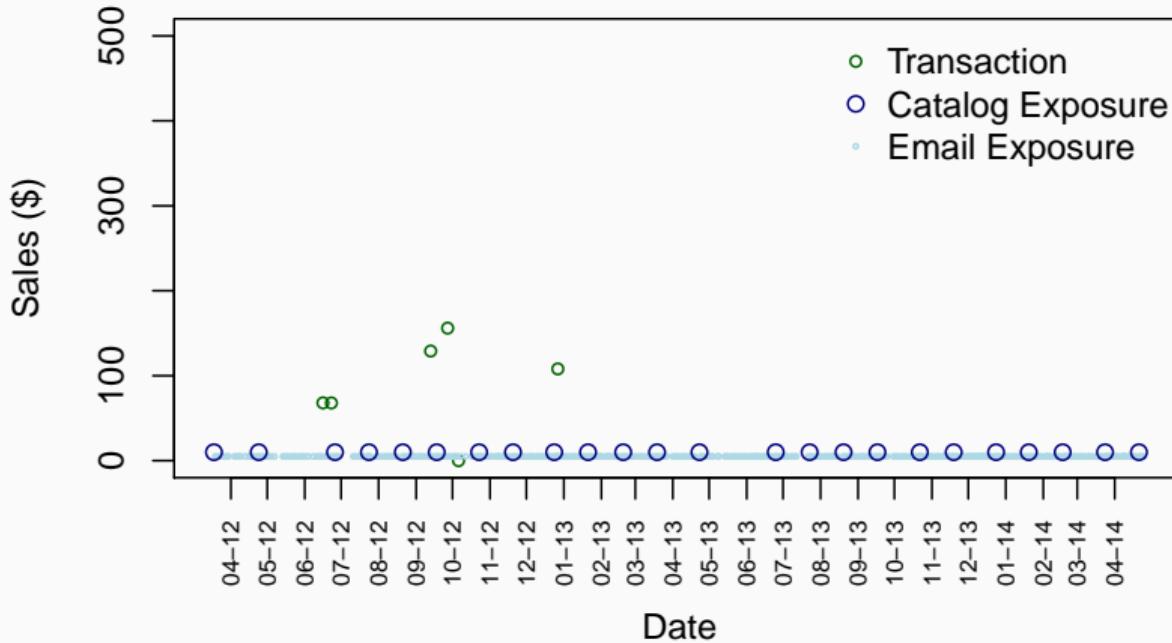
Why not marketing mix modeling?



- As is common in direct marketing, the cadence and reach of campaigns is very steady over time.
- Sales are also quite steady over time.
- Lack of variation makes it difficult to determine any relationship between sales and marketing from aggregate data.

Customer-level advertising response data

Sales for Customer 143



Analyzing campaigns as individual field experiments

Each holdout test can be analyzed individually.

July 2012 Catalog Campaign

	Customers (N)	Ave. 30-Day Sales (\$)	30-Day Purchase Incidence	Ave. Prior 30-Day Sales (\$)
Treated	284	62.34	0.25	58.82
Holdout	16	11.90	0.06	62.65
Difference		50.44	0.18	-3.84
95% CI		(18.95, 81.92)	(0.04, 1)	(-91.82, 84.13)

Attribution model features

Differences in advertising response across customers (heterogeneity)

A tobit model for purchase amount (\$)

Estimated decay rates for catalog and emails

Interaction between email and catalog

No endogeneity correction (We don't need it, as all exposure were randomized in holdout test)

A multi-channel, consumer-level ad-stock model

We relate the weekly total purchase amount, Y_{it} , for each user i in each period t to the advertising exposures, X_{ikt} , on each channel k in all previous periods as follows:

$$Y_{it} = \begin{cases} Y_{it}^* & \text{if } Y_{it}^* > 0 \\ 0 & \text{if } Y_{it}^* \leq 0 \end{cases}$$

$$Y_{it}^* = \mu_i + \sum_k \beta_{ik} W_{ikt} + \sum_{k:k'} \beta_{i,k:k'} W_{i,k:k',t} + \eta_{it}$$

$$W_{ikt} = \sum_{s=0}^{\infty} \rho_{ik}^s (X_{i,k,t-s} + \epsilon_{ikt})$$

$$W_{i,k:k',t} = \sum_{s=0}^{\infty} \rho_{i,k:k'}^s (X_{i,k,t-s} + X_{i,k,t-s} + \epsilon_{ikt})$$

A multi-channel, consumer-level ad-stock model

We relate the weekly total purchase amount, Y_{it} , for each user i in each period t to the advertising exposures, X_{ikt} , on each channel k in all previous periods as follows:

$$Y_{it} = \begin{cases} Y_{it}^* & \text{if } Y_{it}^* > 0 \\ 0 & \text{if } Y_{it}^* \leq 0 \end{cases} \quad \text{Tobit I selection}$$

$$Y_{it}^* = \mu_i + \sum_k \beta_{ik} W_{ikt} + \sum_{k:k'} \beta_{i,k:k'} W_{i,k:k',t} + \eta_{it}$$

$$W_{ikt} = \sum_{s=0}^{\infty} \rho_{ik}^s (X_{i,k,t-s} + \epsilon_{ikt})$$

$$W_{i,k:k',t} = \sum_{s=0}^{\infty} \rho_{i,k:k'}^s (X_{i,k,t-s} + X_{i,k,t-s} + \epsilon_{ikt})$$

A multi-channel, consumer-level ad-stock model

We relate the weekly total purchase amount, Y_{it} , for each user i in each period t to the advertising exposures, X_{ikt} , on each channel k in all previous periods as follows:

$$Y_{it} = \begin{cases} Y_{it}^* & \text{if } Y_{it}^* > 0 \\ 0 & \text{if } Y_{it}^* \leq 0 \end{cases}$$

$$Y_{it}^* = \mu_i + \sum_k \beta_{ik} W_{ikt} + \sum_{k:k'} \beta_{i,k:k'} W_{i,k:k',t} + \eta_{it}$$

Channel-specific stocks

$$W_{ikt} = \sum_{s=0}^{\infty} \rho_{ik}^s (X_{i,k,t-s} + \epsilon_{ikt})$$

$$W_{i,k:k',t} = \sum_{s=0}^{\infty} \rho_{i,k:k'}^s (X_{i,k,t-s} + X_{i,k,t-s} + \epsilon_{ikt})$$

A multi-channel, consumer-level ad-stock model

We relate the weekly total purchase amount, Y_{it} , for each user i in each period t to the advertising exposures, X_{ikt} , on each channel k in all previous periods as follows:

$$Y_{it} = \begin{cases} Y_{it}^* & \text{if } Y_{it}^* > 0 \\ 0 & \text{if } Y_{it}^* \leq 0 \end{cases}$$

$$Y_{it}^* = \mu_i + \sum_k \beta_{ik} W_{ikt} + \sum_{k:k'} \beta_{i,k:k'} W_{i,k:k',t} + \eta_{it}$$

$$W_{ikt} = \sum_{s=0}^{\infty} \rho_{ik}^s (X_{i,k,t-s} + \epsilon_{ikt})$$

Interactions

$$W_{i,k:k',t} = \sum_{s=0}^{\infty} \rho_{i,k:k'}^s (X_{i,k,t-s} + X_{i,k,t-s} + \epsilon_{ikt})$$

A multi-channel, consumer-level ad-stock model

We relate the weekly total purchase amount, Y_{it} , for each user i in each period t to the advertising exposures, X_{ikt} , on each channel k in all previous periods as follows:

$$Y_{it} = \begin{cases} Y_{it}^* & \text{if } Y_{it}^* > 0 \\ 0 & \text{if } Y_{it}^* \leq 0 \end{cases}$$

$$Y_{it}^* = \mu_i + \sum_k \beta_{ik} W_{ikt} + \sum_{k:k'} \beta_{i,k:k'} W_{i,k:k',t} + \eta_{it}$$

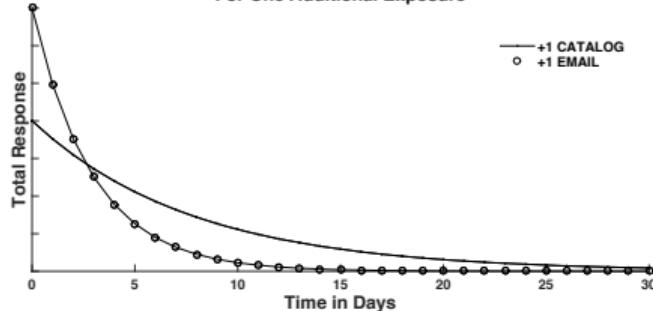
$$W_{ikt} = \sum_{s=0}^{\infty} \rho_{ik}^s (X_{i,k,t-s} + \epsilon_{ikt})$$

$$W_{i,k:k',t} = \sum_{s=0}^{\infty} \rho_{i,k:k'}^s (X_{i,k,t-s} + X_{i,k,t-s} + \epsilon_{ikt})$$

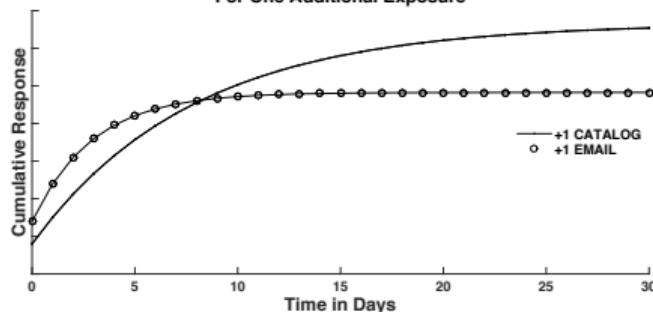
Heterogeneity

Catalog versus email impulse response

Panel A: Comparison Between Cumulative Impulse Responses
For One Additional Exposure

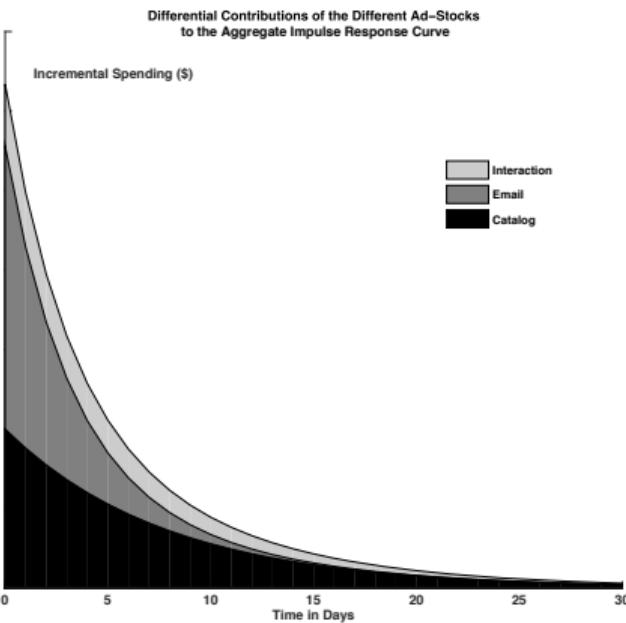


Panel B: Comparison Between Cumulative Impulse Responses
For One Additional Exposure



Simulating the total incremental sales gained by sending each customer a catalog during the first week of the year versus e-mail, we find that catalog is slightly more effective than e-mail.

Combined catalog and email impulse response



Sending both an email and a catalog during the same week produces additional incremental sales lift.

Review

Why not search ads?

Google used jimmy choo shoes

All Shopping Images News Maps More Settings Tools

About 770,000 results (0.82 seconds)

Shop for used jimmy choo shoes on Google

Jimmy Choo Tan Patent ... \$130.00 The RealReal	Jimmy Choo Gold Glitter ... \$171.50 The RealReal	Jimmy Choo Thistle \$505.00 Tradesy Free shipping	Jimmy Choo Metallic ... \$91.00 The RealReal	Jimmy Choo Black Patent ... \$171.50 The RealReal	Jimmy Choo Brown ... \$177.50 The RealReal 32% price drop
---	---	--	--	---	--

Sponsored

Sale - Up to 70% off at Tradesy - tradesy.com

Ad www.tradesy.com/Louis-Vuitton ▾
4.4 ★★★★ rating for tradesy.com
Up to 70% Off Authentic LV! Order Today. Shipping Included.
Authenticity Guaranteed - New Items Added Daily - Rare + Vintage Bags - Buy & Sell Your Clothing

Jimmy Choo on Sale - Up to 70% off at Tradesy
https://www.tradesy.com › Shop ▾
The ultimate destination for guaranteed authentic Jimmy Choo Shoes, Bags & more at up to 70% off.
New and preowned, with safe shipping and friendly returns.

Jimmy Choo Shoes - Up to 70% off at Tradesy
https://www.tradesy.com › Shop › Shoes ▾
The ultimate destination for Jimmy Choo shoes at up to 70% off. New and preowned pumps, heels,
heels, wedges and more, with safe shipping and friendly ...

Jimmy Choo on Sale | Pre-Owned Jimmy Choo | SnobSwap
https://snobswap.com/shop/brand/jimmy-choo ▾
Shop new and pre-owned Jimmy Choo shoes at SnobSwap. Visit our online consignment boutiques for
the best prices on coveted Jimmy Choo pumps, sandals ...

You may have noticed that our example didn't include search advertising and I didn't talk about search except for the search geo-test.

Why not search ads?

The screenshot shows a Google search results page for the query "used jimmy choo shoes". At the top, there's a shopping action bar with options like All, Shopping, Images, News, Maps, More, Settings, and Tools. Below it, a message says "About 770,000 results (0.82 seconds)".

A prominent feature is a "Shop for used jimmy choo shoes on Google" section with six items listed:

Image	Name	Price	Offer	Seller
	Jimmy Choo Tan Patent ...	\$130.00		The RealReal
	Jimmy Choo Gold Glitter ...	\$171.50		The RealReal
	Jimmy Choo Thistle	\$505.00	Trade\$ 91.00	Trade\$ 91.00
	Jimmy Choo Metallic ...	\$91.00	Free shipping	Trade\$ 91.00
	Jimmy Choo Black Patent ...	\$171.50		The RealReal
	Jimmy Choo Brown ...	\$177.50	32% price drop	The RealReal

Below this, there's a sponsored link from Tradesy:

Sale - Up to 70% off at Tradesy - tradesy.com
Ad www.tradesy.com/Louis-Vuitton ▾
4.4 ★★★★ rating for tradesy.com
Up to 70% Off Authentic LV! Order Today. Shipping Included.
Authenticity Guaranteed - New Items Added Daily - Rare + Vintage Bags - Buy & Sell Your Clothing

Organic search results follow:

- Jimmy Choo on Sale - Up to 70% off at Tradesy**
<https://www.tradesy.com> • Shop ▾
The ultimate destination for guaranteed authentic Jimmy Choo Shoes, Bags & more at up to 70% off.
New and preowned, with safe shipping and friendly returns.
- Jimmy Choo Shoes - Up to 70% off at Tradesy**
<https://www.tradesy.com> • Shop • Shoes ▾
The ultimate destination for Jimmy Choo shoes at up to 70% off. New and preowned pumps, heels, heels, wedges and more, with safe shipping and friendly ...
- Jimmy Choo on Sale | Pre-Owned Jimmy Choo | SnobSwap**
<https://snobswap.com/shop/brand/jimmy-choo> ▾
Shop new and pre-owned Jimmy Choo shoes at SnobSwap. Visit our online consignment boutiques for the best prices on coveted Jimmy Choo pumps, sandals ...

You may have noticed that our example didn't include search advertising and I didn't talk about search except for the search geo-test.

Why?

Why not search ads?

The screenshot shows a Google search results page for the query "used jimmy choo shoes". At the top, there's a sponsored shopping ad titled "Shop for used jimmy choo shoes on Google" featuring six Jimmy Choo shoe models with their prices: Tan Patent (\$130.00), Gold Glitter (\$171.50), Thistle (\$505.00), Metallic (\$91.00), Black Patent (\$171.50), and Brown (\$177.50). Below the ad, there's a "Sale - Up to 70% off at Tradesy - tradesy.com" listing with a 4.4-star rating and a link to the website. The main search results include links to Tradesy, Jimmy Choo's official website, and SnobSwap, all advertising discounts on Jimmy Choo shoes.

You may have noticed that our example didn't include search advertising and I didn't talk about search except for the search geo-test.

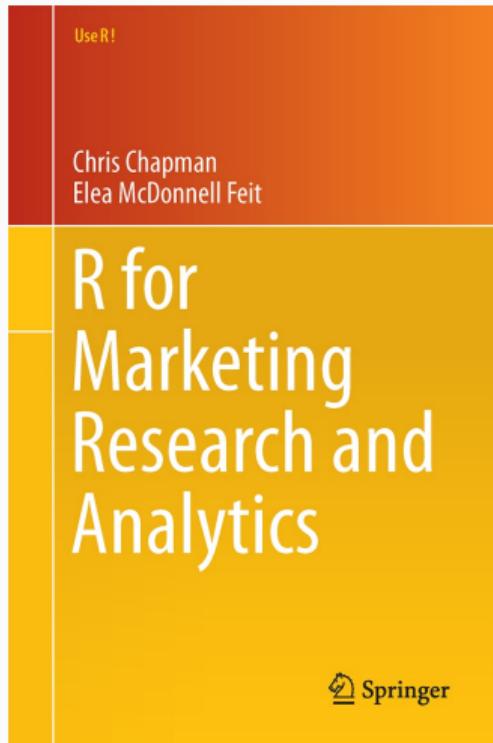
Why?

Google doesn't report search ad impressions for individual users (only clicks). We can include search clicks in our attribution models as an "impression", but they aren't really impressions. There is heavy self-selection in who clicks on a search ad.

Methods for estimating advertising response

Method	Pros	Cons
Attribution rules	Easy to compute	Ignores non-transacting customers Provides no estimate of ad response Sensitive to differences in ad cadence
Marketing-mix modeling	Radically reduces data size	Doesn't work when advertising is steady Can be sensitive to model assumptions Potential endogeneity bias
Model-based attribution	Often works when mix doesn't Can bring user-level controls	Potential endogeneity bias (but less) Large data
Holdout testing	Best causal evidence Simple analysis	Requires planning

If you want to learn more R



An introduction to R for marketing research practitioners that assumes no previous programming experience.

Advanced chapters provide “how-to” guides to segmentation, hierarchical modeling, factor analysis, SEM, market-basket analysis and choice modeling, with some Bayesian applications.

Thank you!

Elea McDonnell Feit

Assistant Professor of Marketing, Drexel University
Senior Fellow, Wharton Customer Analytics Initiative

efeit@drexel.edu

@eleafeit

If you would like to translate this tutorial into another tool such as Python, R/tidyverse or SAS, please let me know. I'd be happy to collaborate with you.



Thank you!

For those who are registered for
the Opening Reception of our
conference, please join us at the
The Logan Hotel at 6:30 PM

#wcaiconf



A Map of
PHILADELPHIA
and Parts Adjacent.

By
N. Scull and G. Heap.

The Distances of particular
Places in this Map;